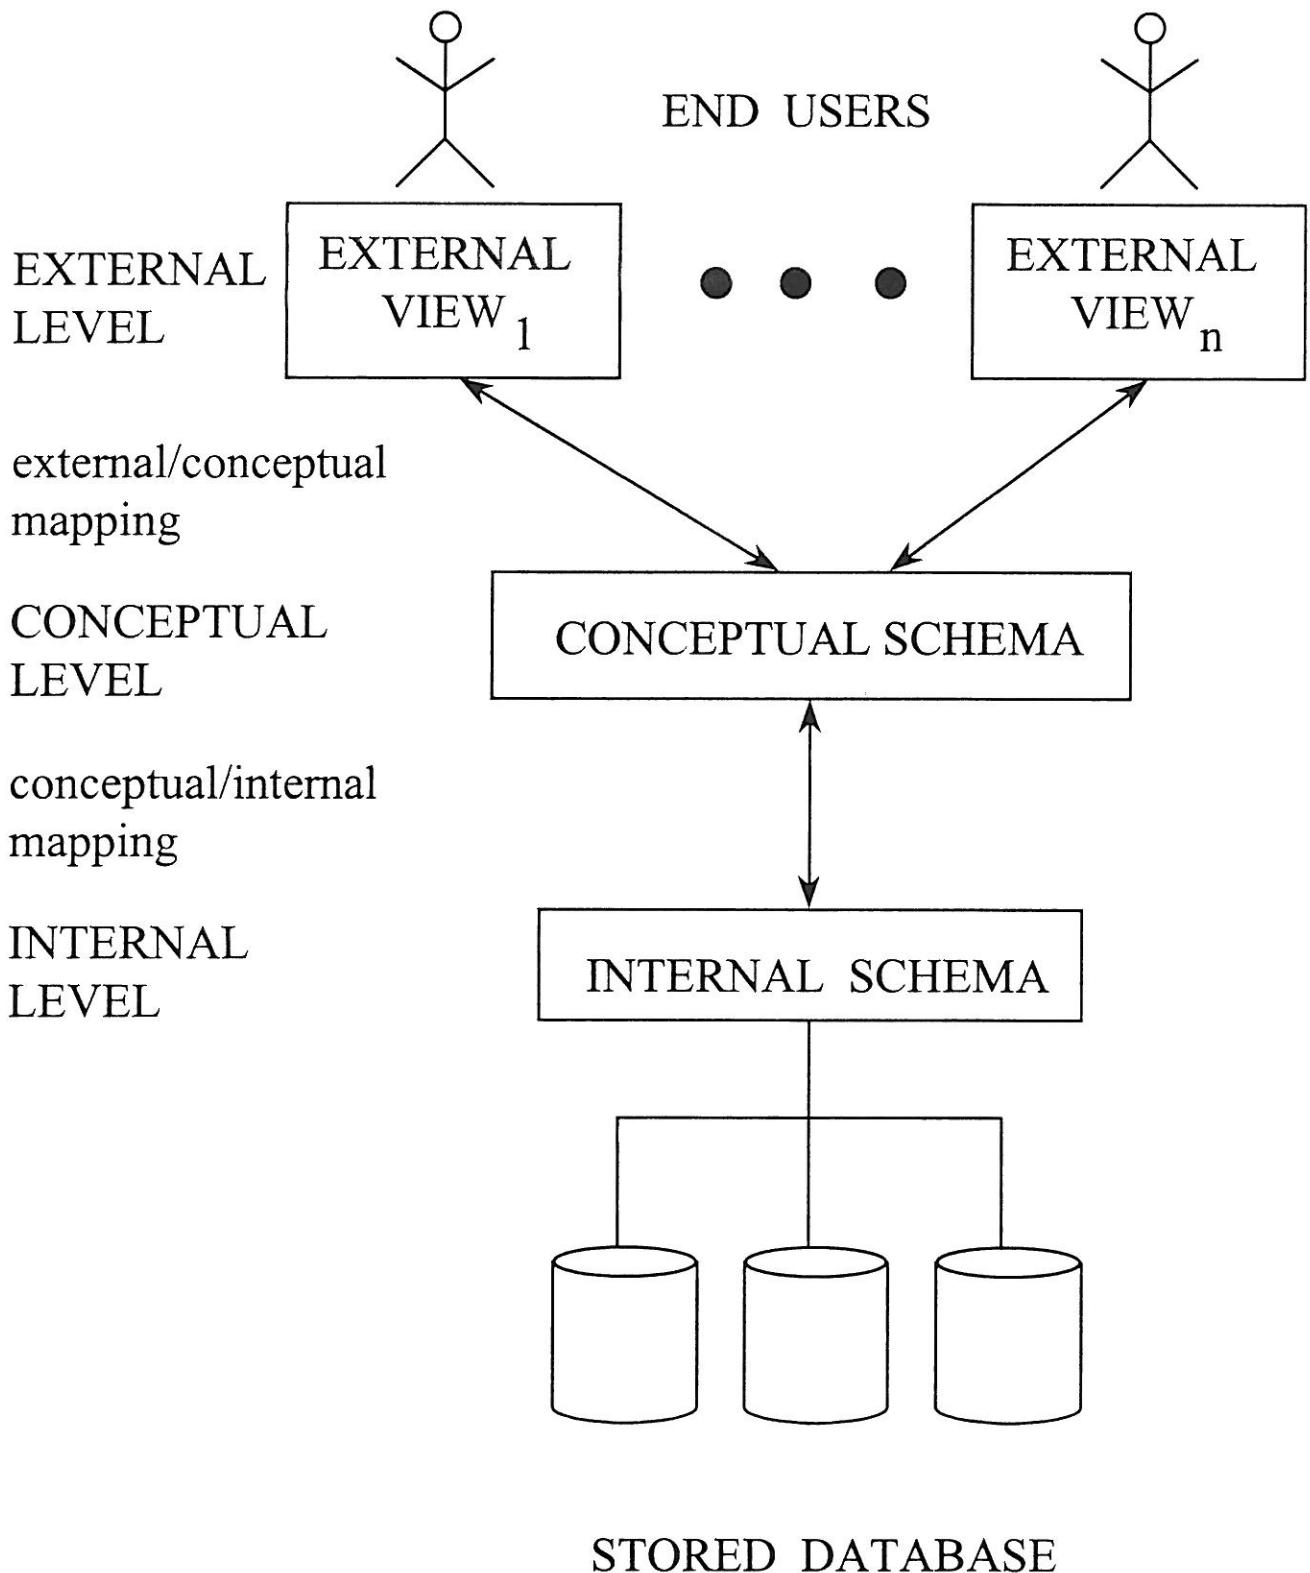


# Introduction to Database Modeling

Peter Scheuermann

Dept. of EECS

[peters@eeecs.northwestern.edu](mailto:peters@eeecs.northwestern.edu)



**ANSI / SPARC ARCHITECTURE**

## **OUTLINE:**

1. DATABASE ARCHITECTURE FRAMEWORK  
file processing versus database approach
2. DATA MODELING WITH THE ENTITY-RELATIONSHIP APPROACH
3. INTRODUCTION TO THE RELATIONAL MODEL
4. TRANSFORMATION OF E-R SCHEMA TO RELATIONAL SCHEMA
5. NORMALIZATION IN THE RELATIONAL MODEL
6. OVERVIEW OF SQL

# THE ENTITY-RELATIONSHIP MODEL

- DESCRIPTION OF THE “REAL WORLD” -  
an important tool in the communication among  
users, system analysts, database administrators

1. **ENTITY:** anything that can be distinctly identified, i.e. has existence in the real world

- **ENTITY SET:** a group of entities with identical properties

EMPLOYEE

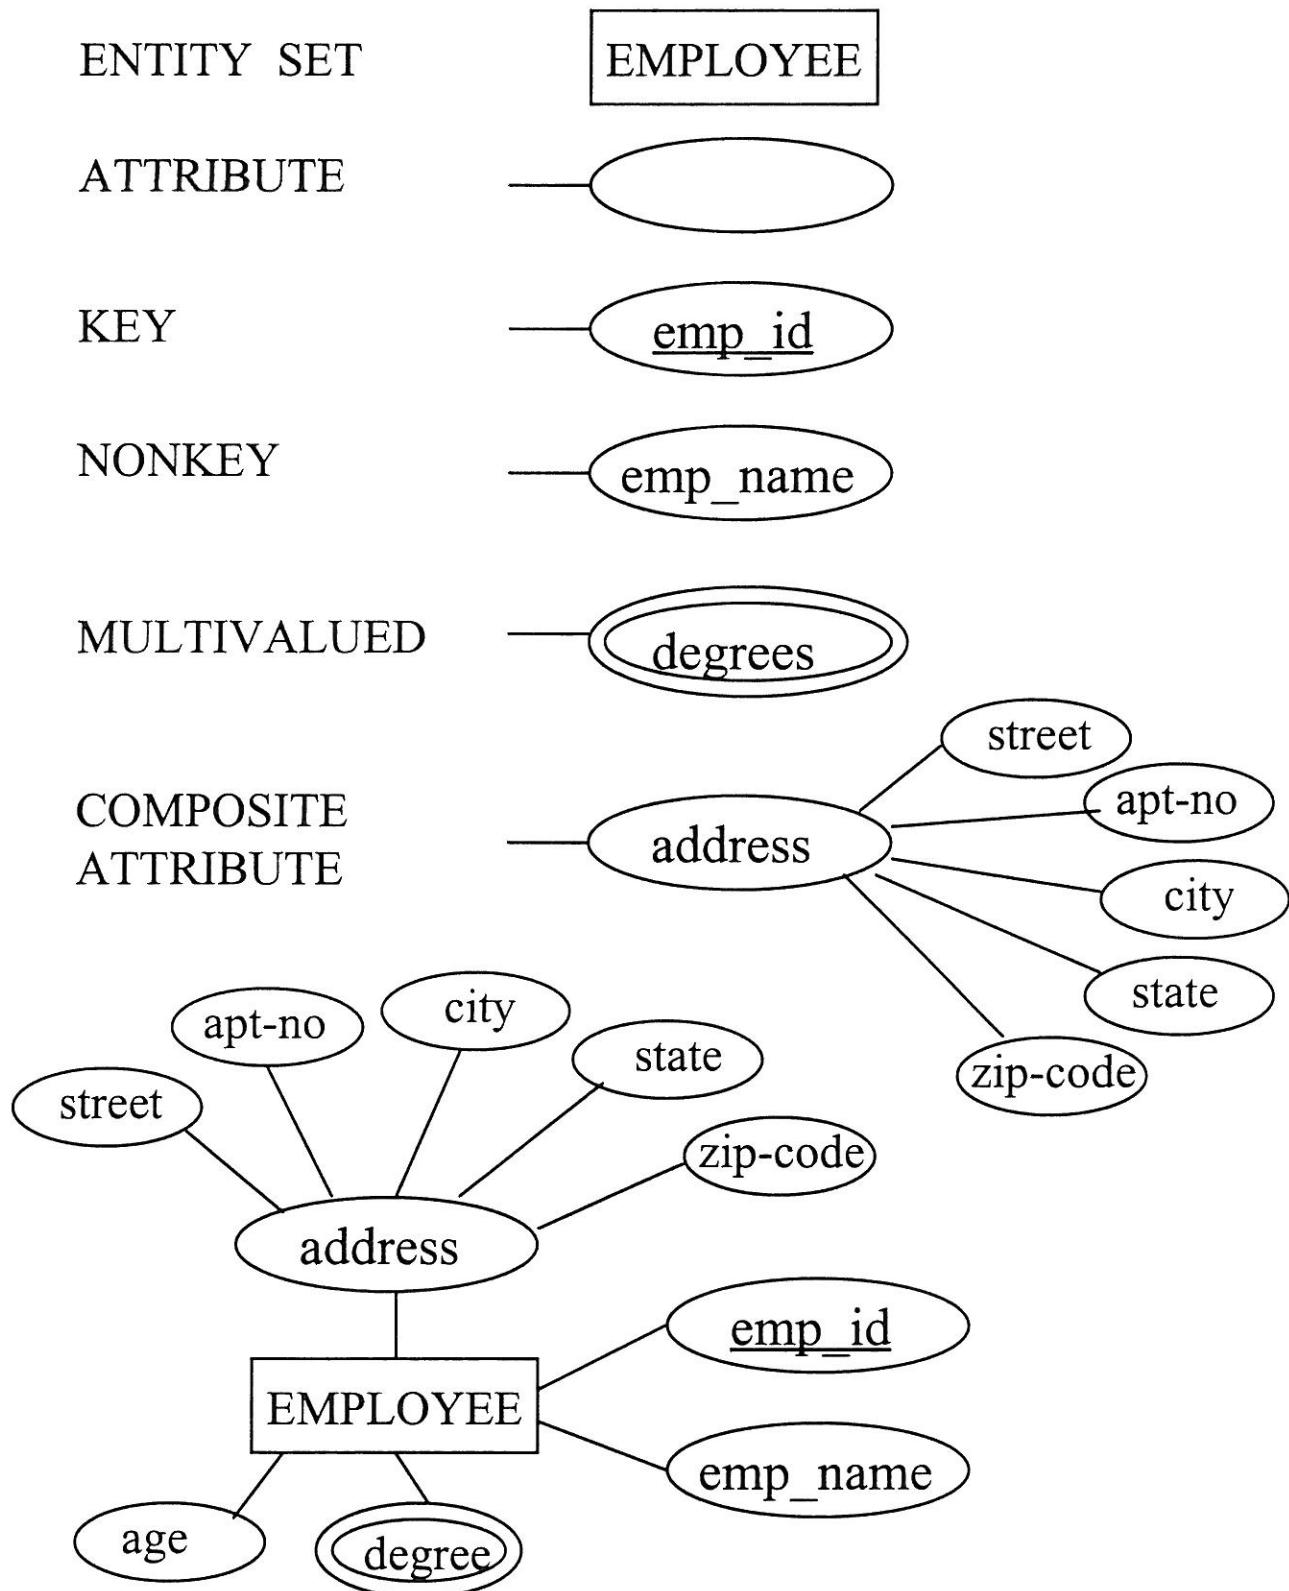
- **DESCRIPTION OF ENTITIES**

- **ATTRIBUTES:** ~ characteristics of attributes that provide descriptive detail

Ex: employee name, age, salary

- each entity attribute can have *single* or *multiple values*

# CONCEPTS AND THEIR REPRESENTATION



- **ENTITY KEY:** an attribute ( or set of attributes) which uniquely identifies an entity within an entity set

- MOST ATTRIBUTES ARE *SINGLE VALUED*

- THE MEANINGS OF *NULL VALUES*:

- (a) **NOT APPLICABLE**

- single family home has no appt-no.

- (b) **UNKNOWN**

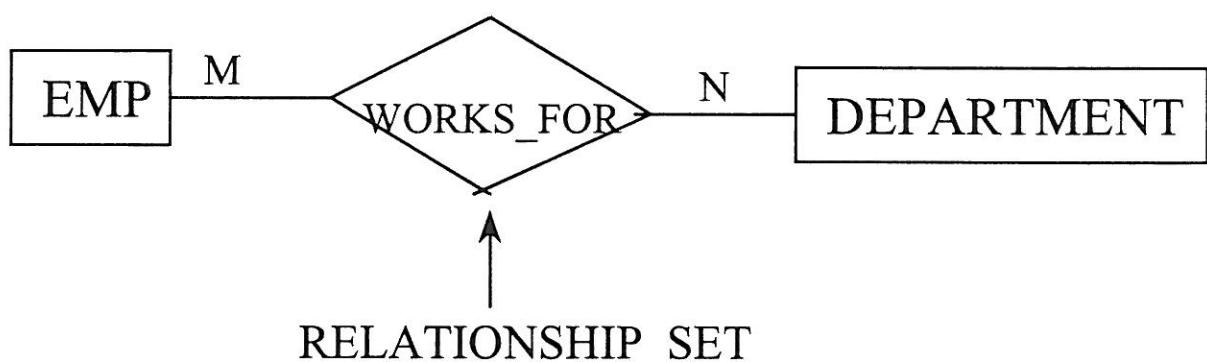
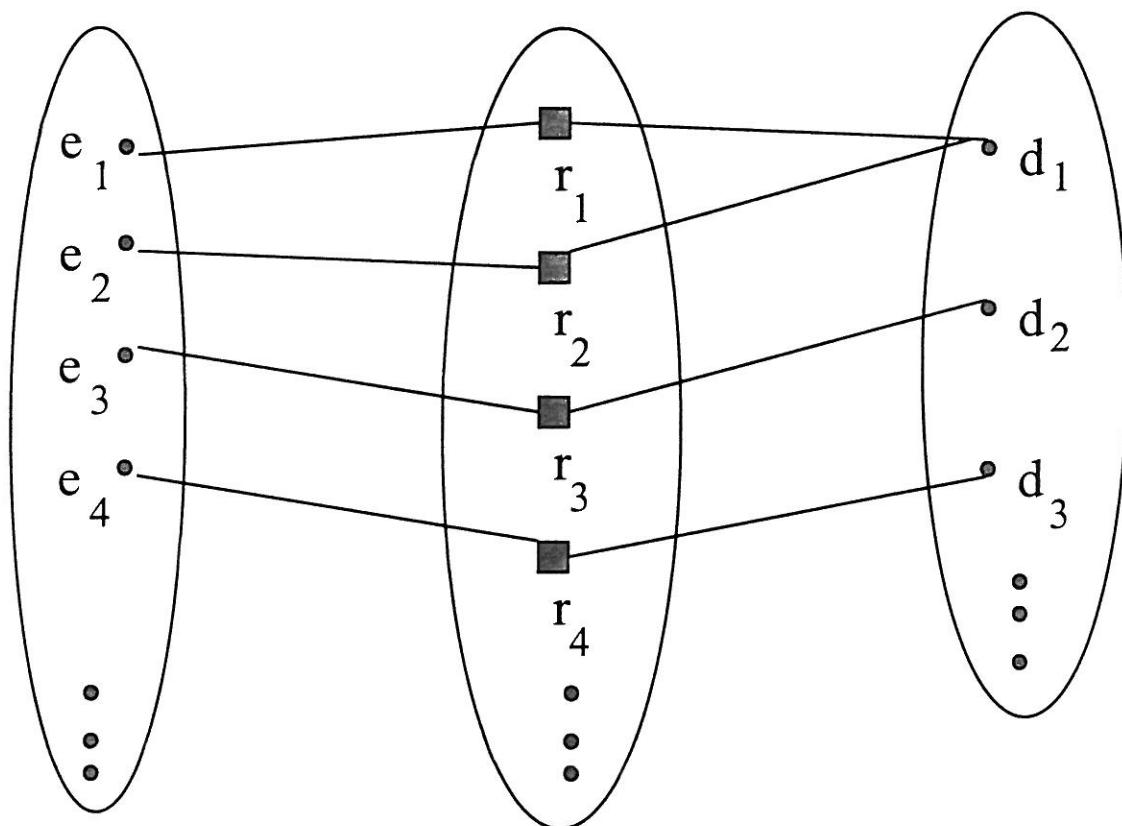
- we don't know the state where Mary Smith lives.

- **VALUE SETS (DOMAINS) :**

- each simple attribute is associated with a value set which specifies the set of values that may be legally assigned to that attribute

Ex: age value set : [16 , 70]

**2. RELATIONSHIP:** association (mapping) among entities where association includes exactly one entity from each participating entity set

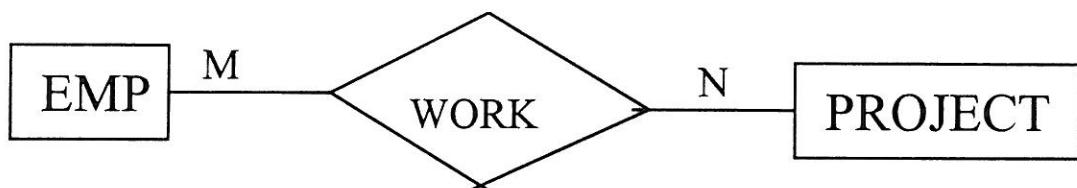


- RELATIONSHIPS ARE CHARACTERIZED IN TERMS OF
  - degree*
  - cardinality*
  - constraints*
  - attributes*

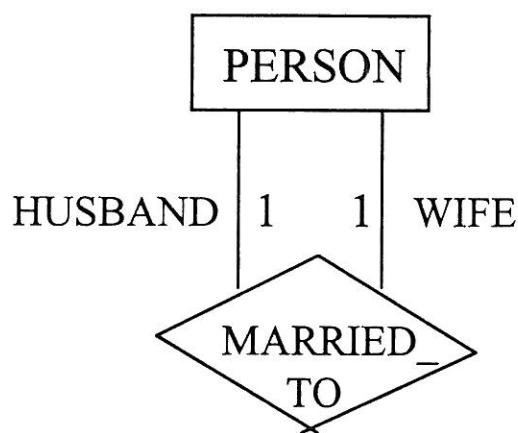
- DEGREE OF RELATIONSHIP SET:

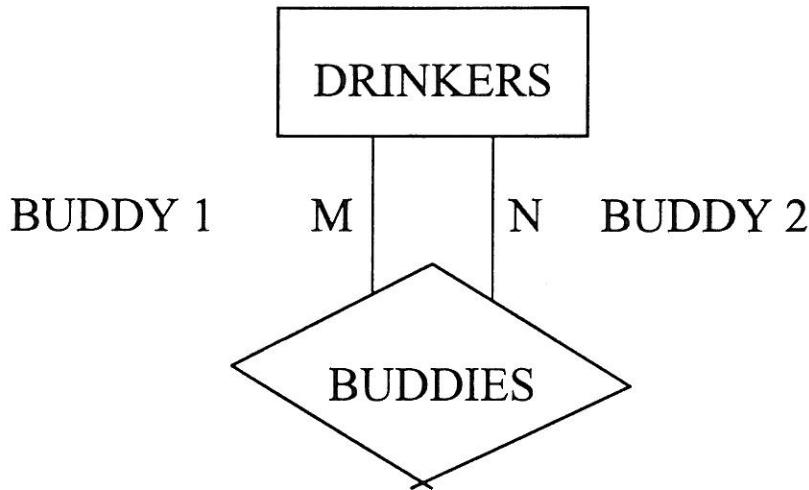
no. of participating entity sets

**BINARY:**



**RECURSIVE BINARY:**





Buddy 1	Buddy 2
$d_1$	$d_2$
$d_1$	$d_3$
$d_2$	$d_1$
$d_2$	$d_4$

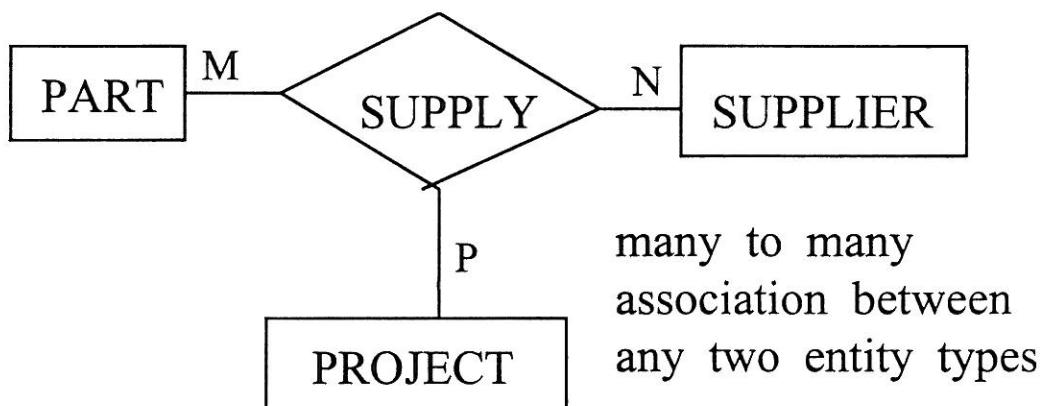
*Buddies* IS symmetric, *married* not

- no way to model “symmetric” in E/R

### Design question:

Should we replace husband and wife by one relationship spouse?

## TERNARY:

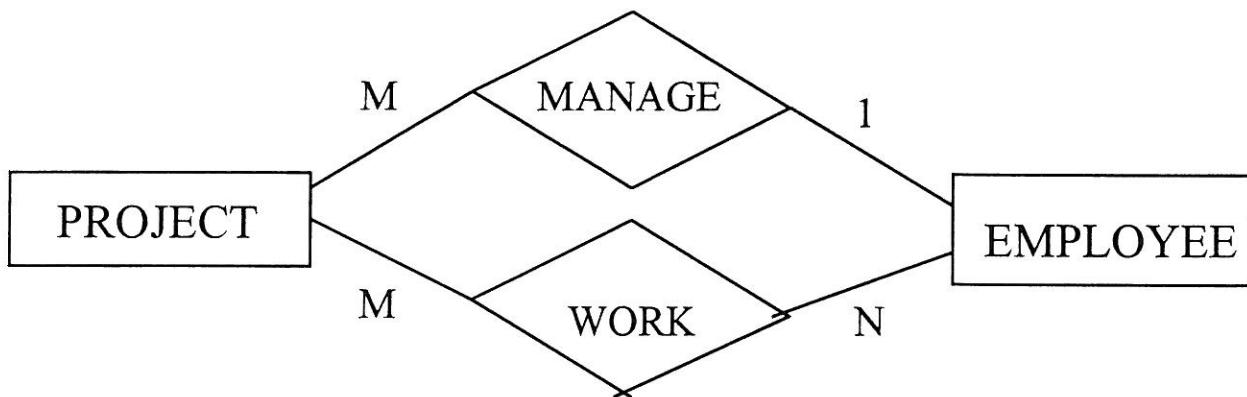


## RELATIONSHIP SET TABLE:

PART#	SUPPLIER#	PROJ#
25	4	1
25	5	2
10	4	2
17	4	3
17	2	1
17	5	1

- *ternary* relationships (and higher) should be avoided if possible
- but *ternary* relationship is necessary when two (or three) *binary* relationships are not equivalent

- **CARDINALITY RATIO:** specifies the number of relationship instances that an entity can participate in

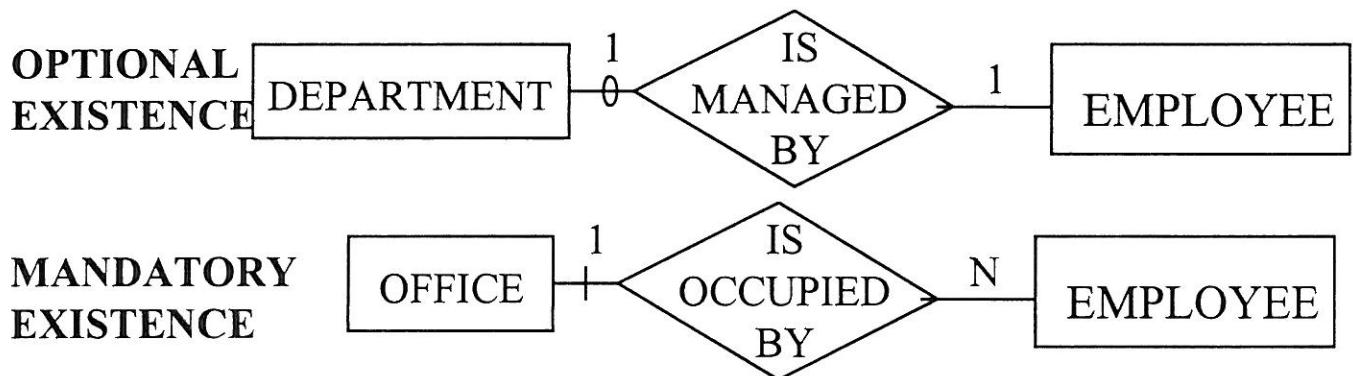


POSSIBLE MAPPINGS ARE 1:1 , M : 1 , 1 : M , M : N

- each project has only *one* manager , but an employee can manage *many* projects

**IMPORTANT:** *one* and *many* refers to maximum cardinality

## •CONSTRAINTS ON PARTICIPATION:



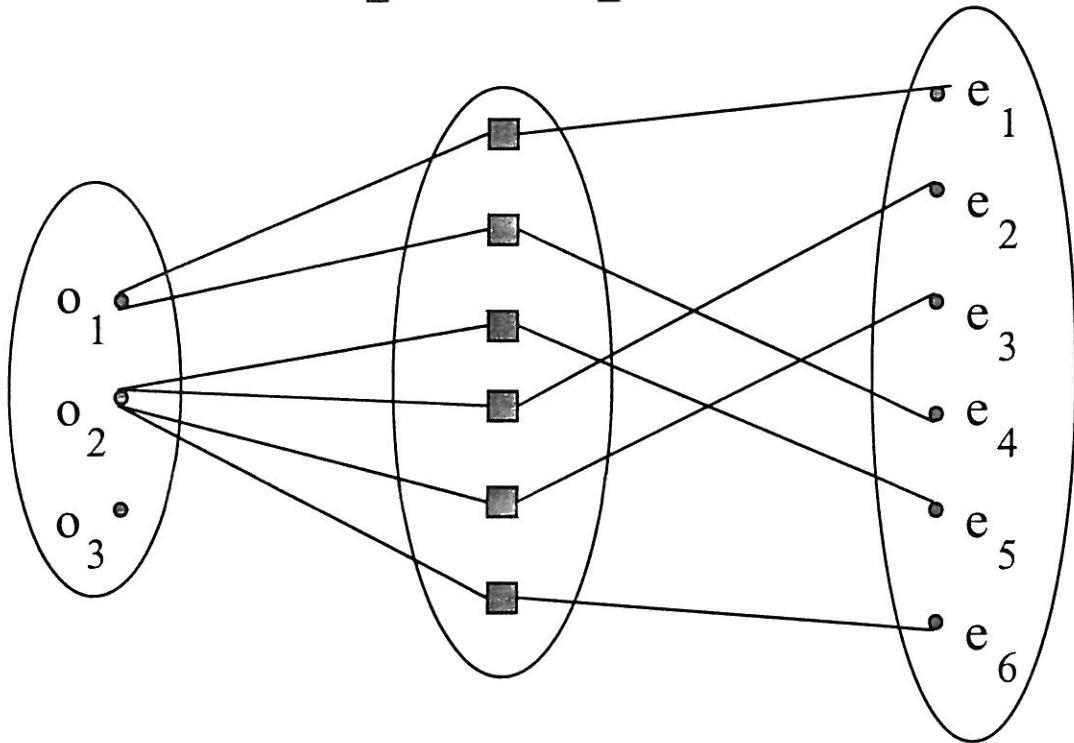
**OPTIONAL EXISTENCE:** defines a minimum cardinality of zero

**MANDATORY EXISTENCE:** defines a minimum cardinality of one

OFFICE

IS\_OCCUPIED\_BY

EMPLOYEE

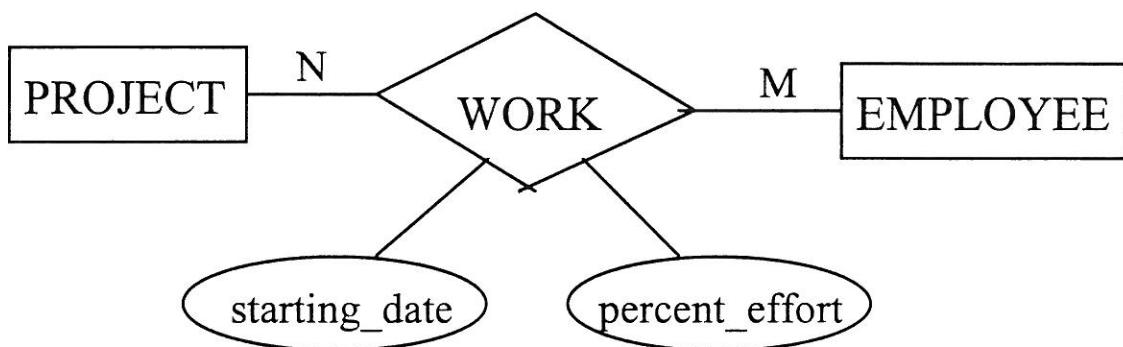


OFFICE HAS *MANDATORY* EXISTENCE IN IS\_OCCUPIED\_BY



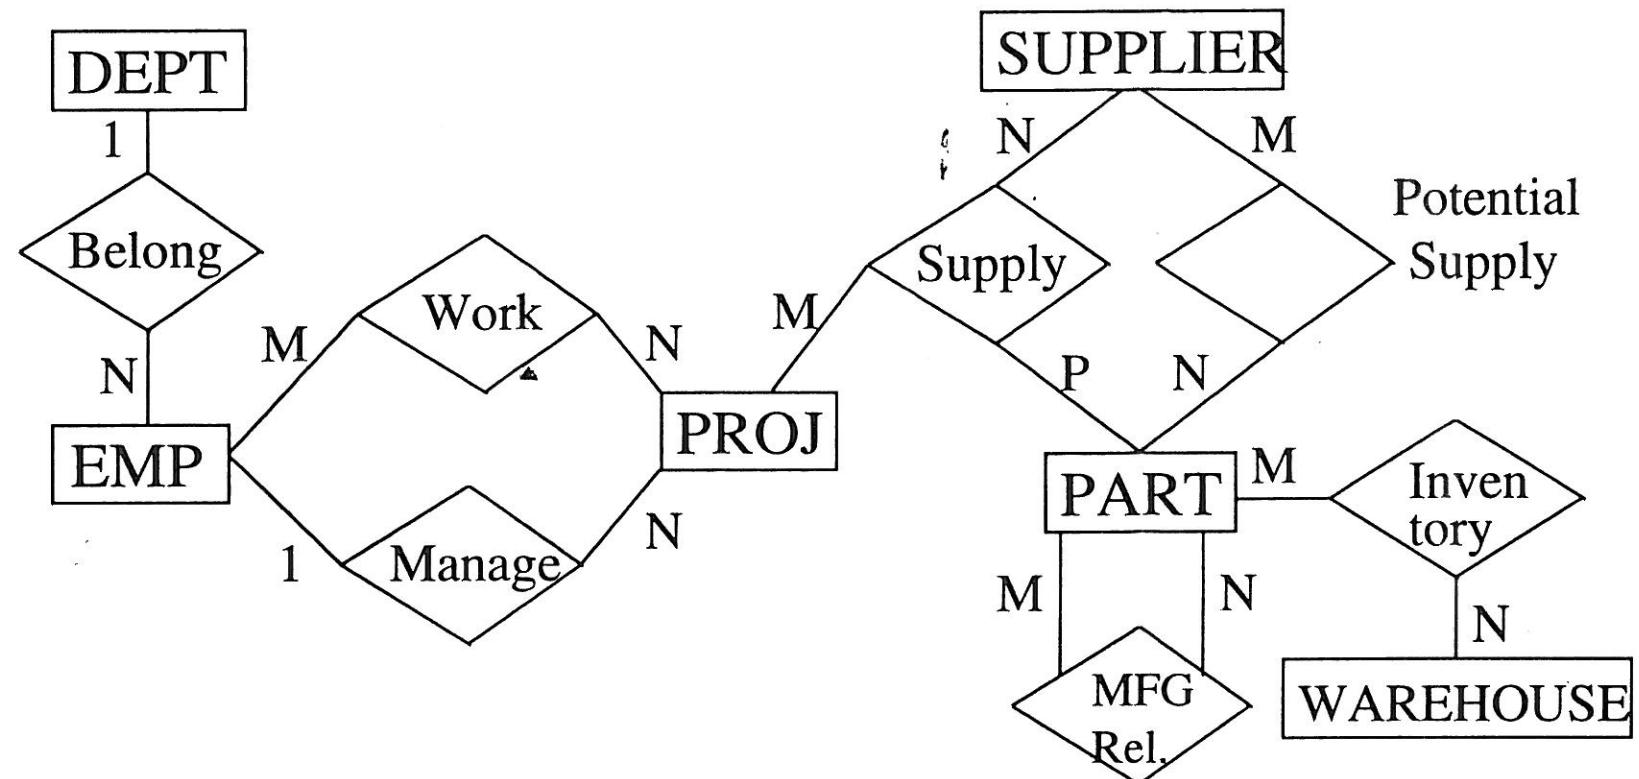
every employee is assigned exactly one office

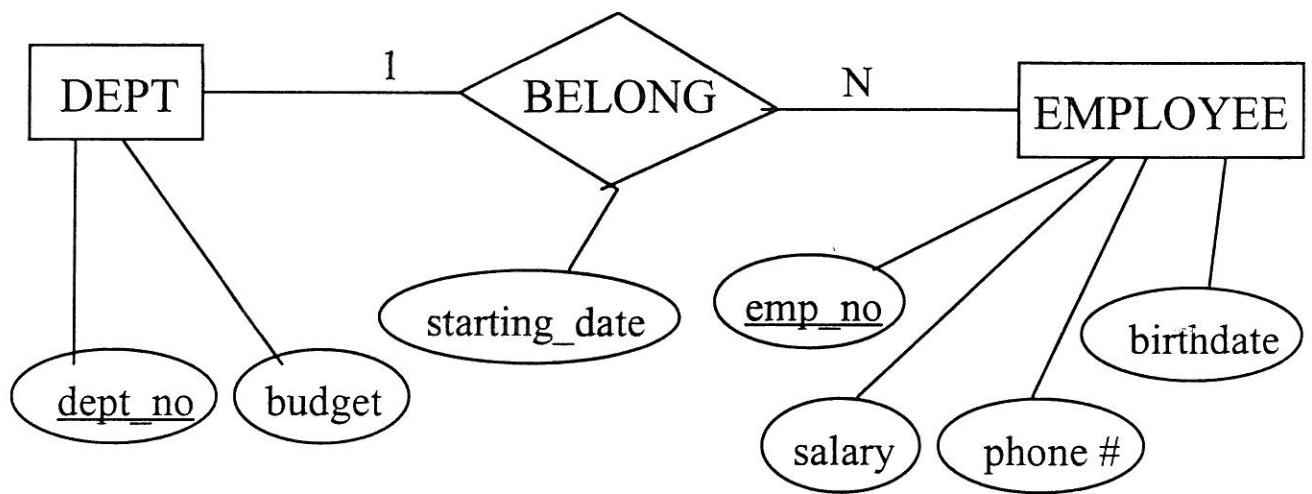
## •RELATIONSHIPS HAVE ATTRIBUTES TOO



# THE ENTITY - RELATIONSHIP MODEL

Example : Manufacturing Company





**NOTE:** attributes *birthdate* and *starting\_date* are defined on the same value-set

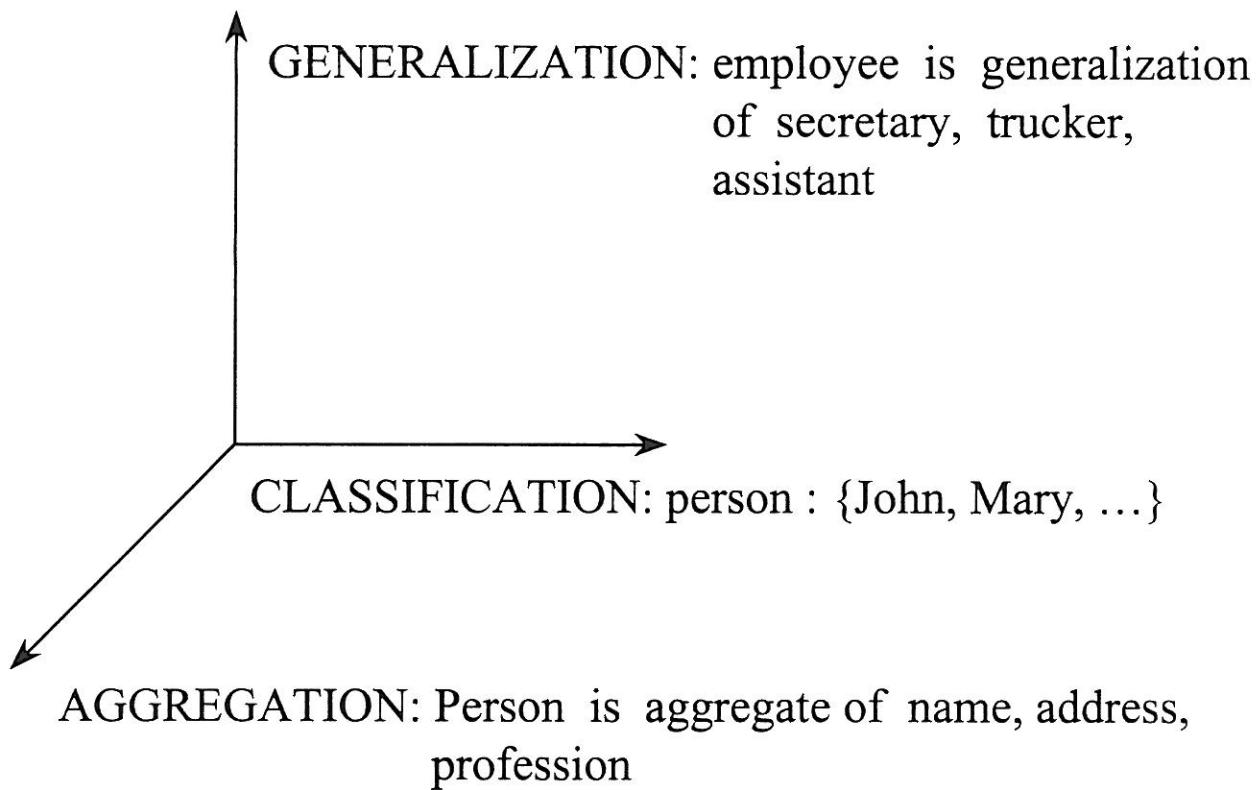
## •ADVANCED CONCEPTS

### •MODELING STRUCTURAL ABSTRACTIONS

(A) **CLASSIFICATION:** collection of objects with similar properties viewed as a class

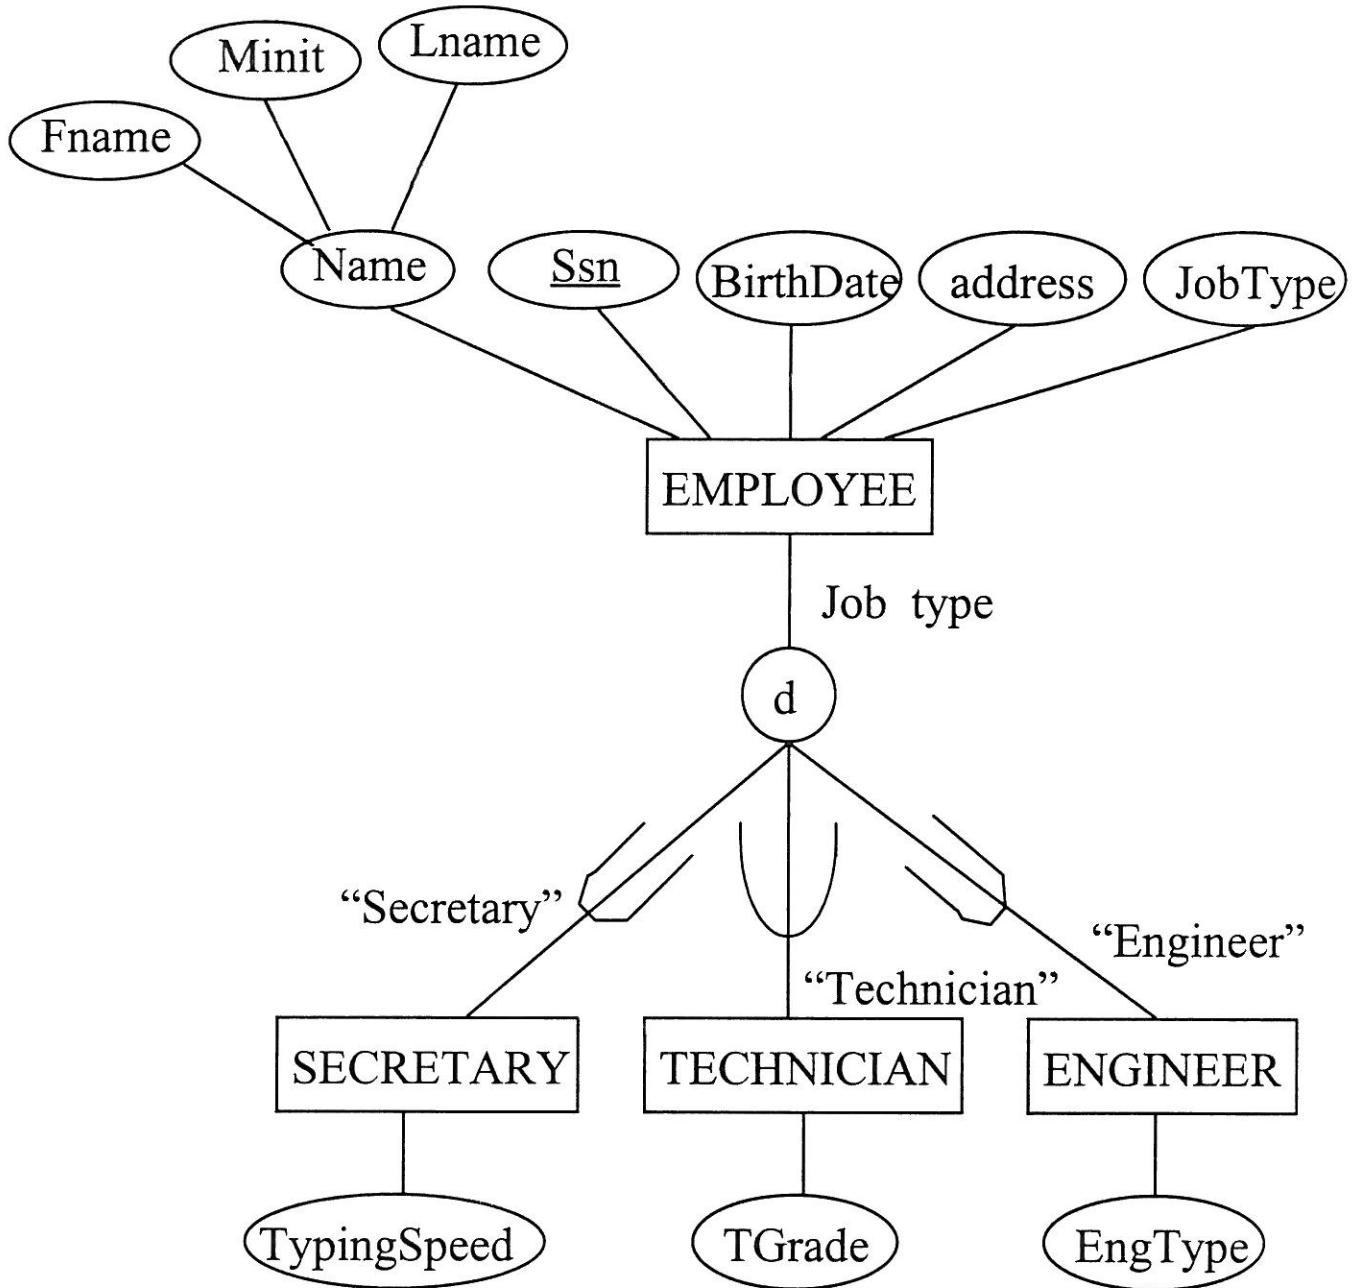
(B) **AGGREGATION:** abstraction by which an object is constructed from its constituent objects

(C) **GENERALIZATION:** set of dissimilar objects is regarded as a single generic object



## GENERALIZATION HIERARCHY

(class/subclass hierarchy)



Note: subset symbol on a line indicates direction of class/subclass relationship.

employee - generalized class contains only attributes in common

secretary, engineer, trucker - specializations of employee

#### •**ATTRIBUTE INHERITANCE**

an entity that is a member of a subclass inherits all the attributes of the entity as a member of the superclass. the entity will also inherit all *relationship* instances for relationship types in which the superclass participates.

#### •**IMPLIED CONSTRAINTS**

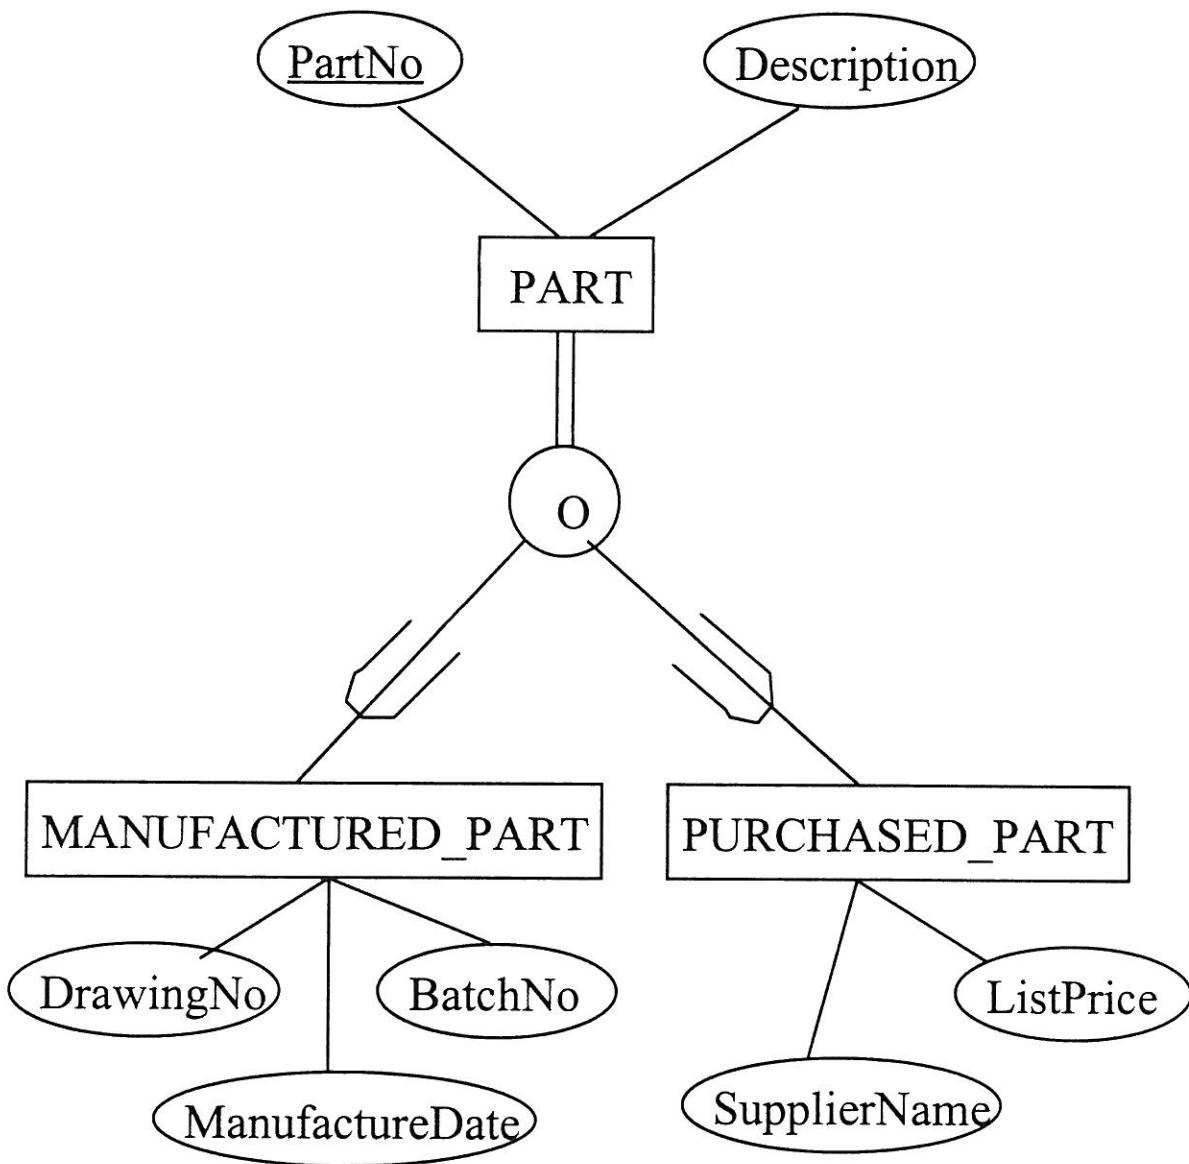
an entity cannot exist in database by being only a member of a subclass, it must also be member of superclass (reverse is not true).

## **ADDITIONAL CONSTRAINTS ON SPECIALIZATION**

### **(A) DISJOINTNESS CONSTRAINT:**

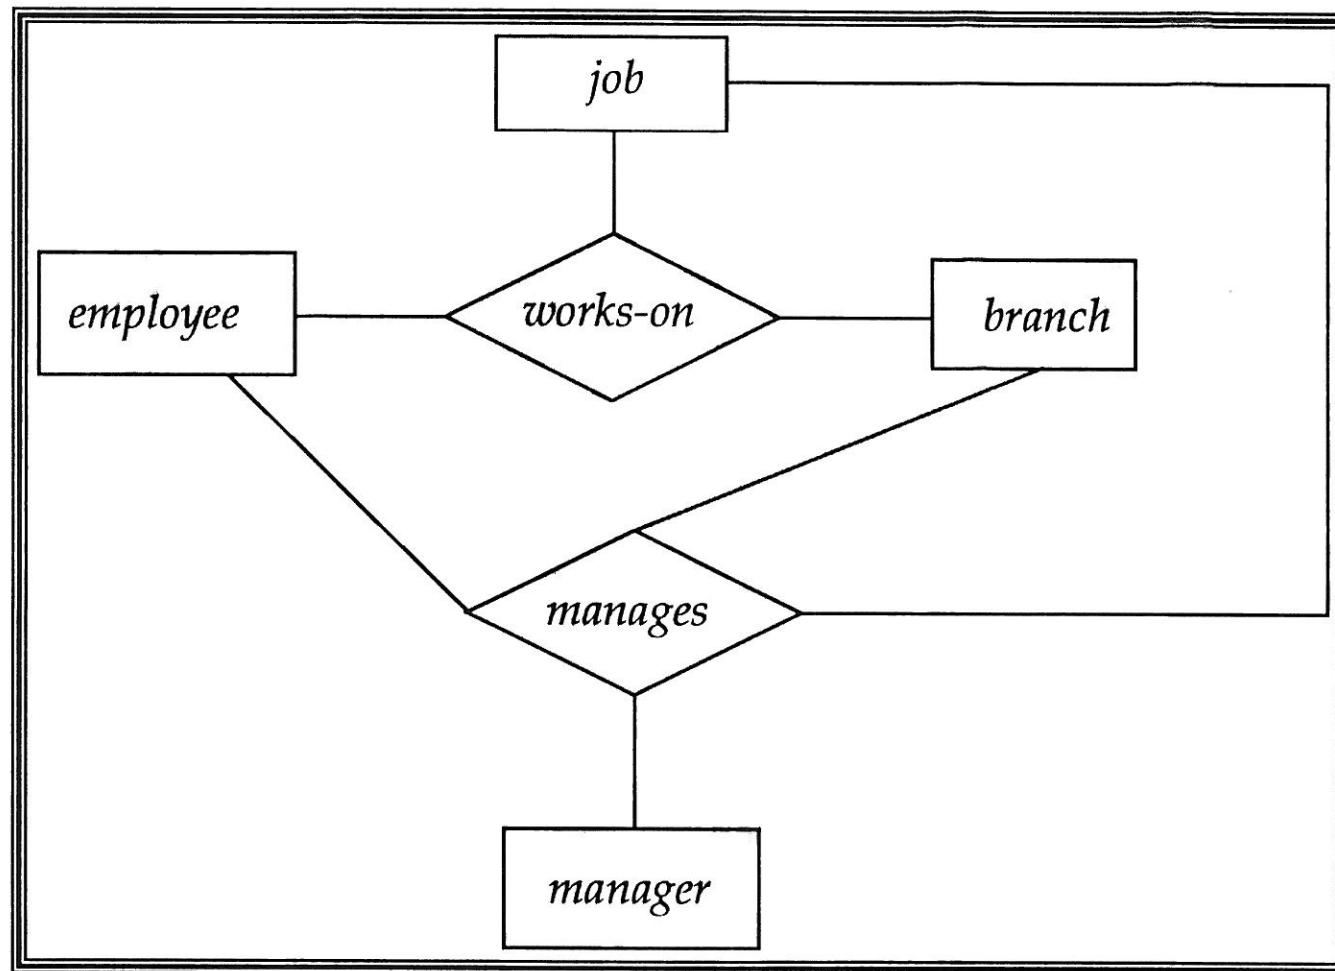
Subclasses of specialization are disjoint

### **(B) OVERLAPPING SUBCLASSES**

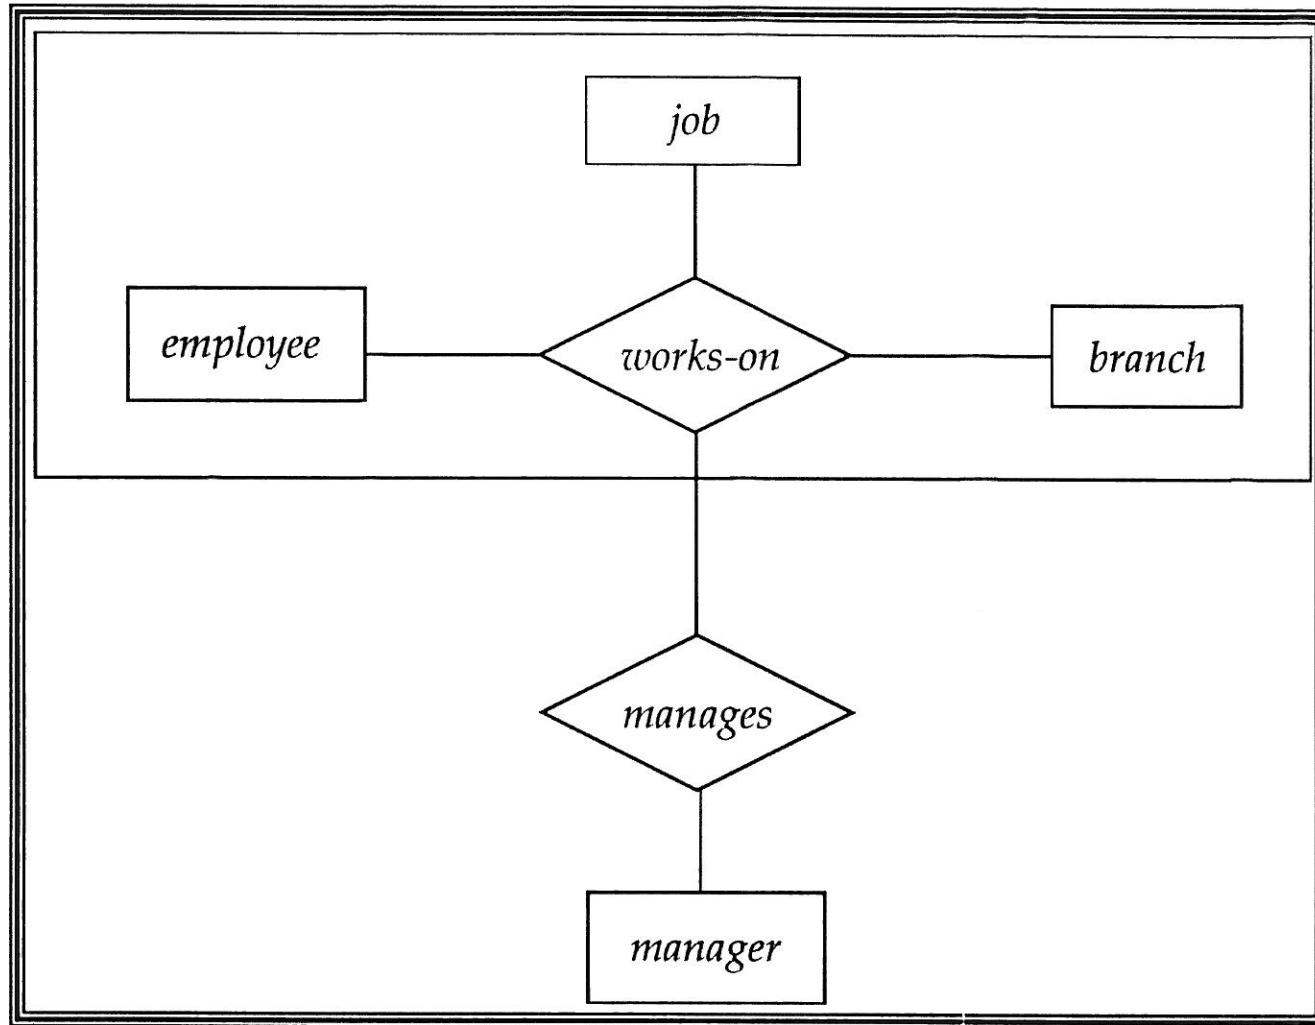


# Aggregation

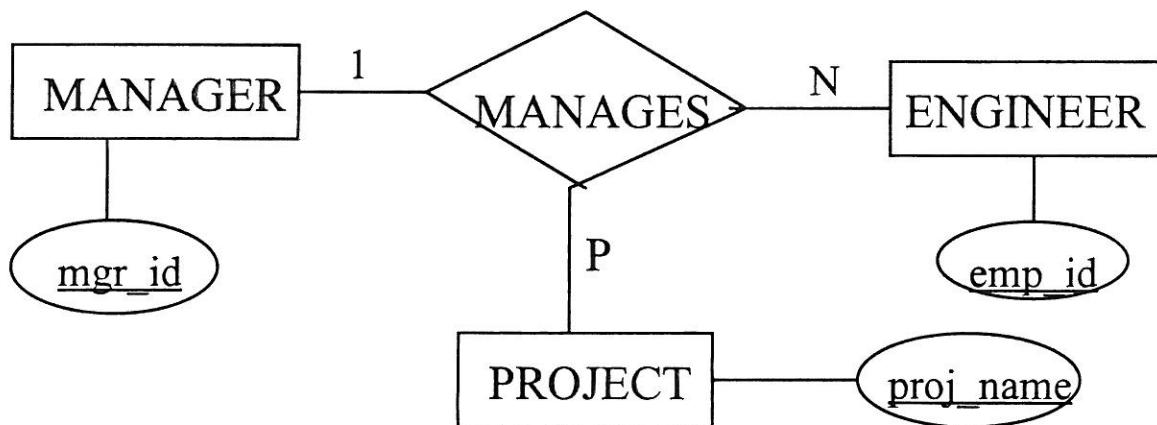
- Consider the ternary relationship *works-on*
- Suppose we want to record managers for tasks performed by an employee at a branch



# E-R Diagram With Aggregation



## DETERMINING THE CARDINALITY OF TERNARY RELATIONSHIPS



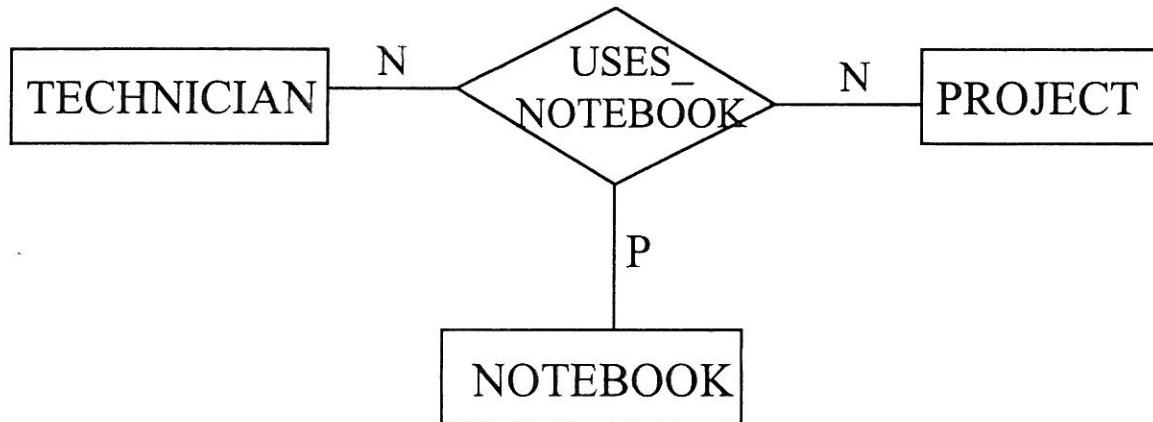
- **ASSERTION 1 :** one engineer, working under one manager could be working on many projects
- **ASSERTION 2 :** one project, under the direction of one manager, could have many engineers
- **ASSERTION 3 :** one engineer, working on one project, must have only a single manager

*EQUIVALENT FUNCTIONAL DEPENDENCY*

$$emp\_id, proj\_name \rightarrow mgr\_id$$

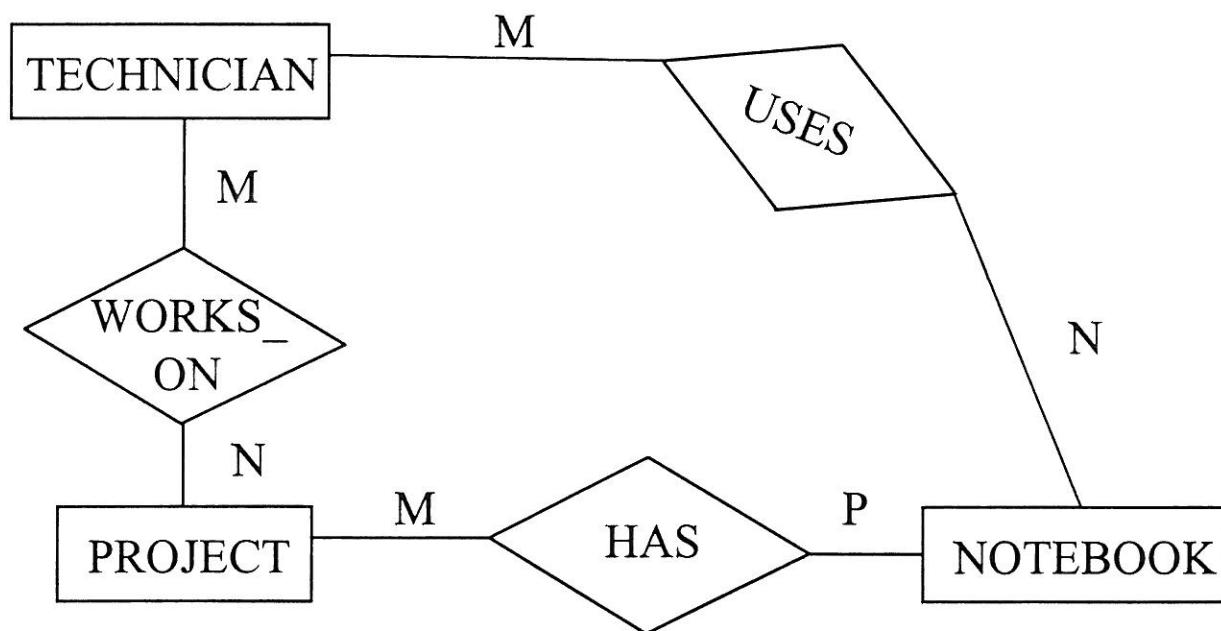
↑  
keys of associated entities

## CONVERTING TERNARY RELATIONSHIPS

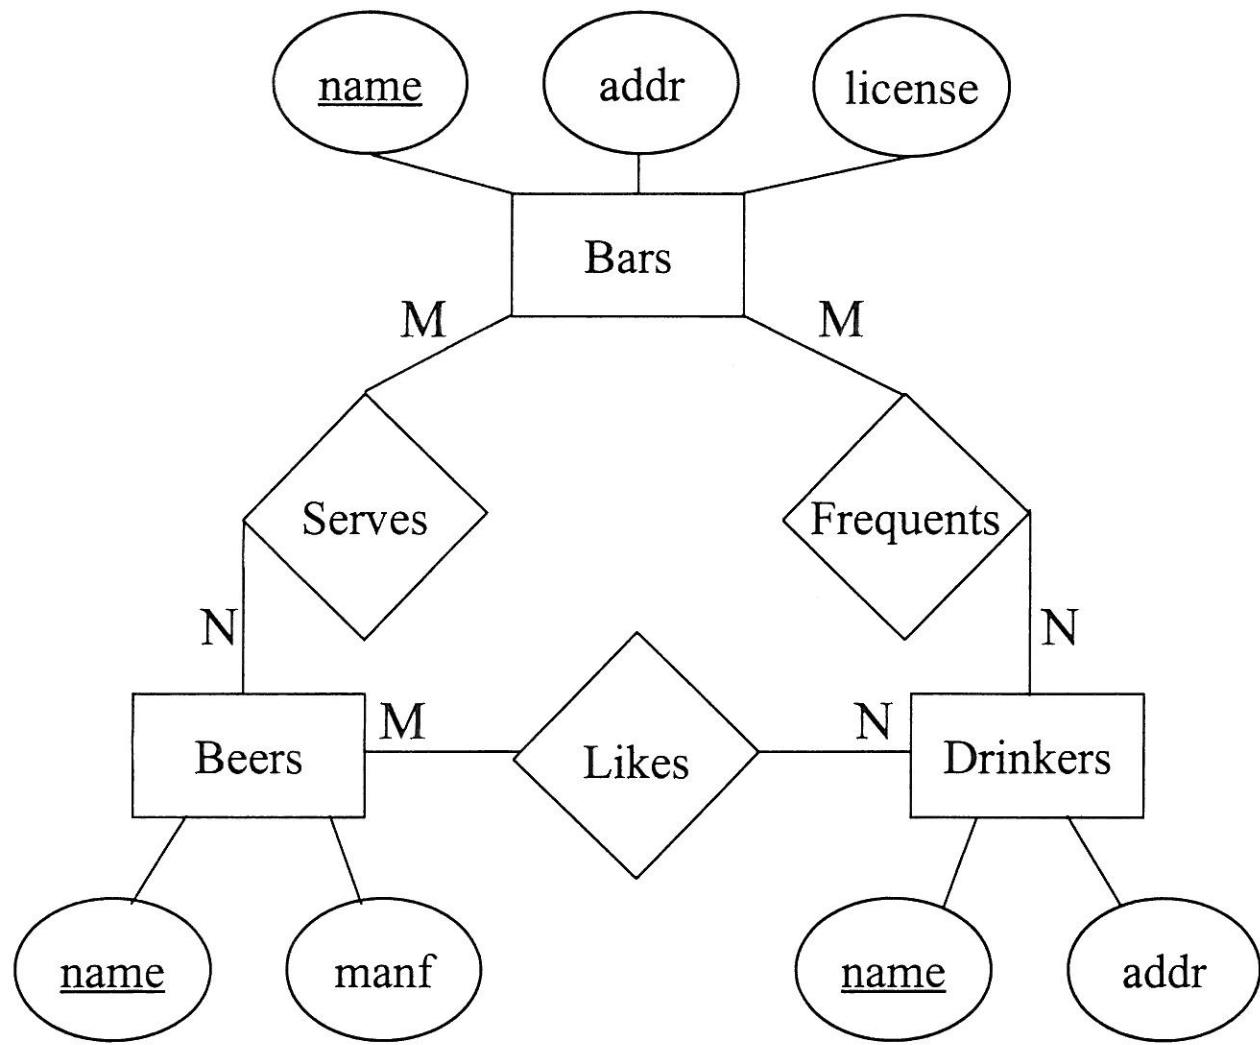


if each technician can be working on several projects and uses the same notebooks on each project;

=> can *decompose* 3-ary relationship into binary relationships

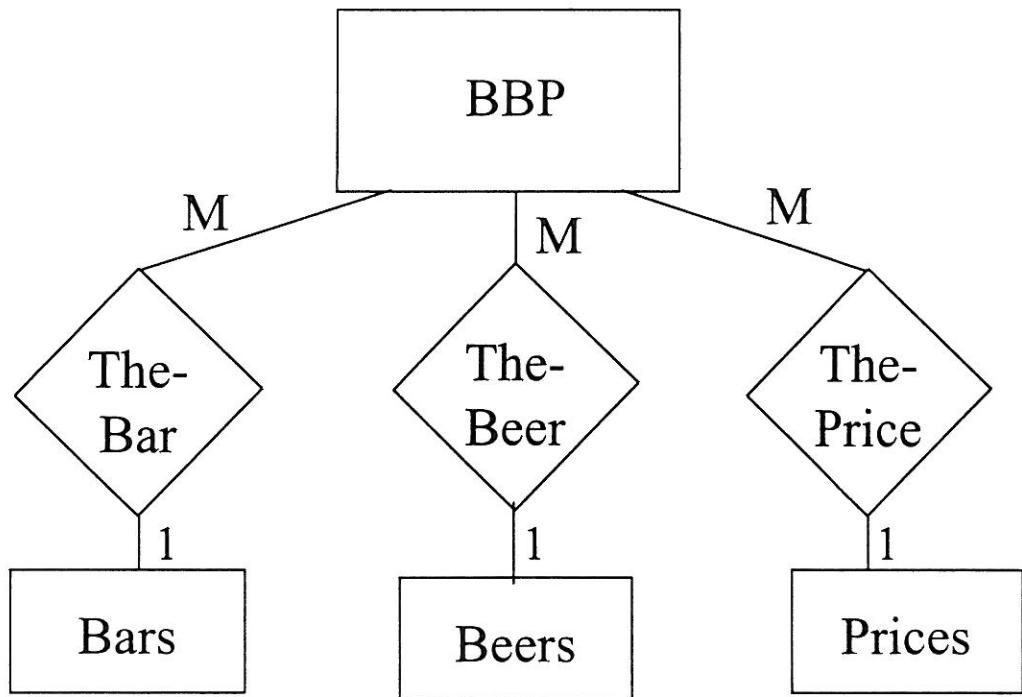


# Beers-Bars-Drinkers Example



## Converting Multiway to 2-Way

- Awkward in E/R, but necessary in certain “object-oriented” models.
- Create a new *connecting* E.S. to represent rows of a relationship set.
  - e.g. (Joe’s Bar, Bud, \$2.50) for the *Sells* relationship.
- Many-one relationships from the connecting E.S. to the others.



## Weak Entity Sets

- Entity set E has only *partial key*  
Its set is formed by combining partial key  
with key(s) of *identifying owner(s)*

*Conditions:*

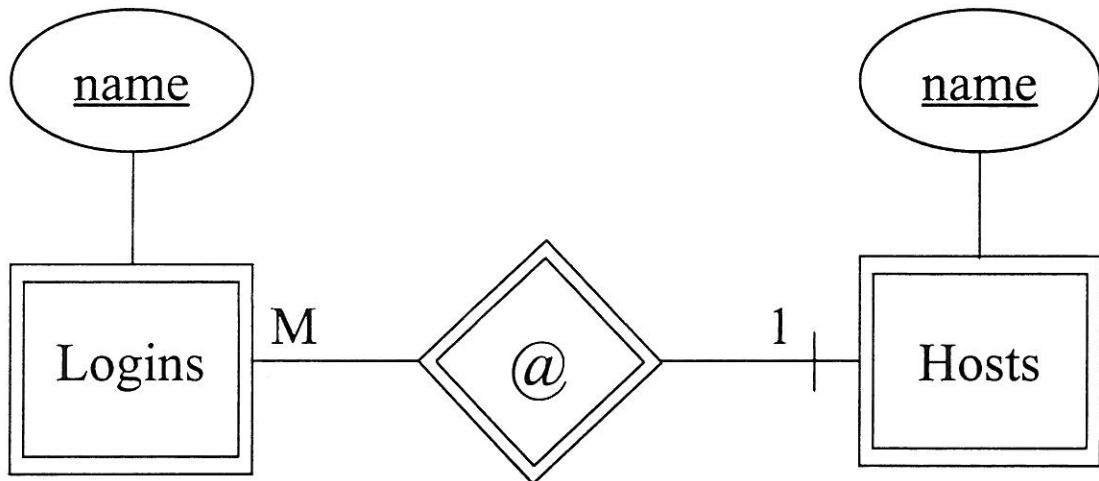
- Owner entity set (s) + weak entity set E  
must participate in 1 : M relationship set
- Identifying owner has *mandatory* existence  
in identifying relationship set

**NOTE:** If identifying owner is itself weak,  
then the key of E is formed recursively

## Example: Logins (Email Addresses)

Login name = user name + host name, e.g.

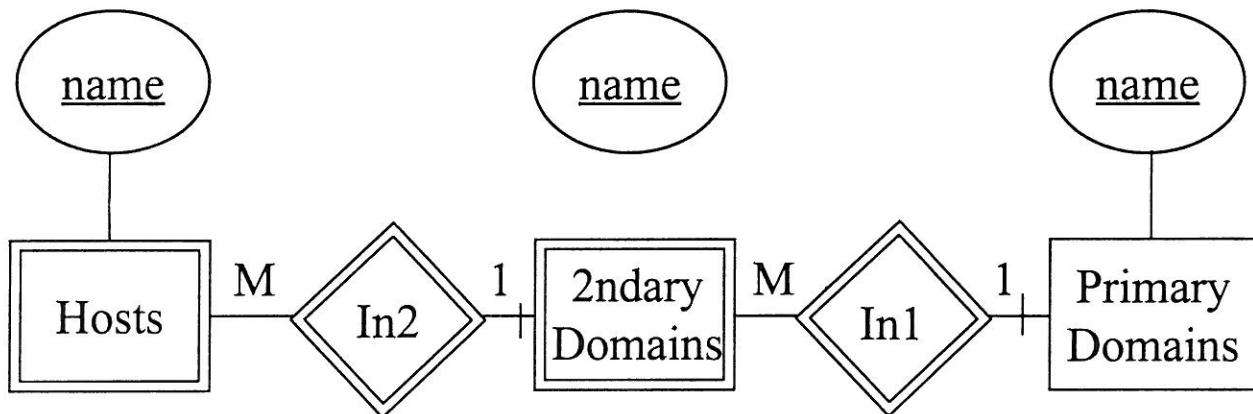
- A “login” entity corresponds to a user name on a particular host, but the password table doesn’t record the host, just the user name, e,g, . peters.
- Key for a login = the user name at the host (which is unique for that host only) + the IP address of the host (which is unique globally) .



- Design Issue: Under what circumstances could we simply make login-name and host- name be attributes of logins, and dispense with the weak E.S.?

## Example: Chain of “Weakness”

Consider IP addresses consisting of a primary domain (e.g. edu) subdomain (e.g. northwestern) and host (e.g. begonia).



- Key for primary domain = its name.
- Key for secondary domain = its name + name of primary domain
- Key for host = its name + key of secondary domain = its name + name of secondary domain + name of primary domain

## **STEPS OF LOGICAL DESIGN WITH E-R MODEL**

- (1) CLASSIFY ENTITIES AND ATTRIBUTES**
  - entities should contain descriptive information
  - classify multivalued attributes as entities
- (2) IDENTIFY GENERALIZATION HIERARCHIES**
- (3) DEFINE RELATIONSHIPS**  
check for redundant relationships
- (4) PERFORM *VIEW INTEGRATION* AND  
CONSTRUCT ONE CONCEPTUAL SCHEMA**
- (5) TRANSFORM E-R MODEL INTO SQL TABLES**
- (6) NORMALIZE THE TABLES**

# THE RELATIONAL MODEL

RELATION => TABLE => SET

- no two tuples are identical
- ordering of tuples is immaterial

RELATION SCHEMA



DEPT (DNAME, THIS-YEAR-BUDGET, LAST-YEAR-BUDGET)

D.Name	This-year-budget	Last-year-budget
ENG.	200 K	100 K
SALES	800 K	600 K
MANUF	500 K	300 K

relation instance

tuple

EMP (EMP-NO, SALARY, OFF-PHONE, HOME-PHONE,  
BIRTHDATE)

Emp-no	Salary	Off.Phone	Home Phone	Birthdate
101	60 K	1-6181	328-0152	1952
117	45 K	1-3298	675-2191	1965
250	90 K	1-5276	321-0712	1943
390	35 K	1-9913	275-3172	1961

BELONG (EMP-NO, DEPT-NAME, START-DATE)

Emp-no	Dept-name	Start-date
101	ENG	1985
117	SALES	1990
250	ENG	1970
390	MANUF	1992

- ALL ATTRIBUTE VALUES ARE ATOMIC

### SUPPLY (SUPPLIER-NO, PART-NO)

Supplier-no	Part-no
S2	P1
S2	P2
S4	P2
S4	P4
S4	P5

correct !

Supplier-no	Part-no
S2	{P1,P2}
S4	{P2,P4,P5}

values are sets  
(repeating groups)

wrong !

**DEFINITION:** a relation is in *first normal form* (1nf) if it contains only atomic values, that is there are no repeating groups within a tuple

- RELATION SCHEMA  $R(A_1, A_2, \dots, A_n)$

- ATTRIBUTE  $A_i$  : role played by domain  $D_i$

- DEGREE of relation is the number of attributes (n) in its relation schema

## •KEYS

- SUPERKEY:** a set of one or more attributes which taken collectively allows us to uniquely identify a tuple in a relation
- KEY:** a *minimal superkey* that is a superkey from which we cannot remove any attributes and still have the uniqueness hold

**Ex :** *EMP'(EMP-NO, EMP-NAME, SALARY, BIRTH-DATE)*

has a number of *Candidate Keys*

- PRIMARY KEY:** candidate key chosen by the system (user) and used to identify tuples in the relation  
(usually *index* also constructed on *primary key*)

•**FOREIGN KEY (FK)** : attribute  $FK$  (possibly composite) of relation  $R_2$  is a foreign key if and only if it satisfies :

1. Each value of  $FK$  is wholly null or wholly non-null
2. There exists a relation  $R_1$  with primary key  $PK$  such that each non-null value of  $FK$  is identical to the value of  $PK$  in some tuple of  $R_1$ .

**Notes :**

- relations  $R_1$  and  $R_2$  need not be distinct
- $R_2$  is called the *referencing relation*
- $R_1$  is called the *referenced relation*

**Ex :**

In relation belong {emp\_no, dept\_name} are two FKs.

# **What makes a system relational (Codd)**

=> Two main integrity rules:

**A. Entity integrity rule:** no component of the primary key of a relation is allowed to accept NULLs.

(remember NULL means: “value unknown”  
or “does not apply”)

**Reason:** if two or more tuples had nulls for primary keys,  
we are not able to distinguish between them

**B. Referential integrity rule:** the database must not contain any unmatched foreign key values.

foreign key & referential integrity rule are defined in terms of one another.

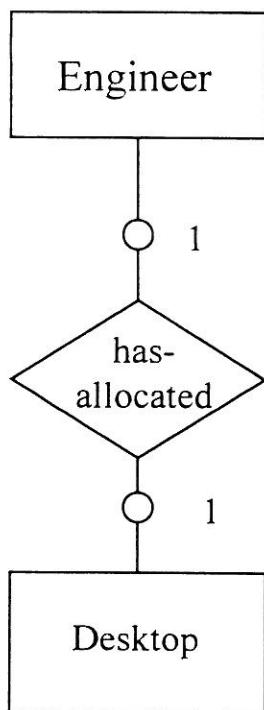
**Important:** these rules do not say how the system should enforce them!

# Transformation rules: From ER model to SQL tables

1. An entity set becomes all table with the same attributes
2. Binary relationship set rules:

## 2.1 One-to-one

relationship set **embedded** in either entity table and the primary key of one entity table appears as **foreign key** in the other.

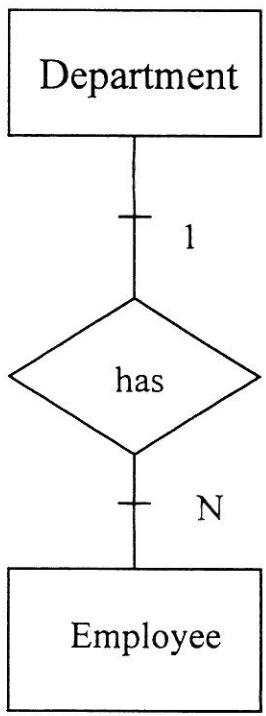


```
create table engineer
  (emp_id char(10),
  desktop_no integer,
  primary key (emp_id),
  foreign key (desktop_no) references
    desktop
  on delete set null on update cascade);
```

```
create table desktop
  (desktop_no integer,
  emp_id char(10),
  primary key (desktop_no),
  foreign key (emp_id) references
    engineer on delete set null
  on update cascade);
```

## 2.2 One-to-many

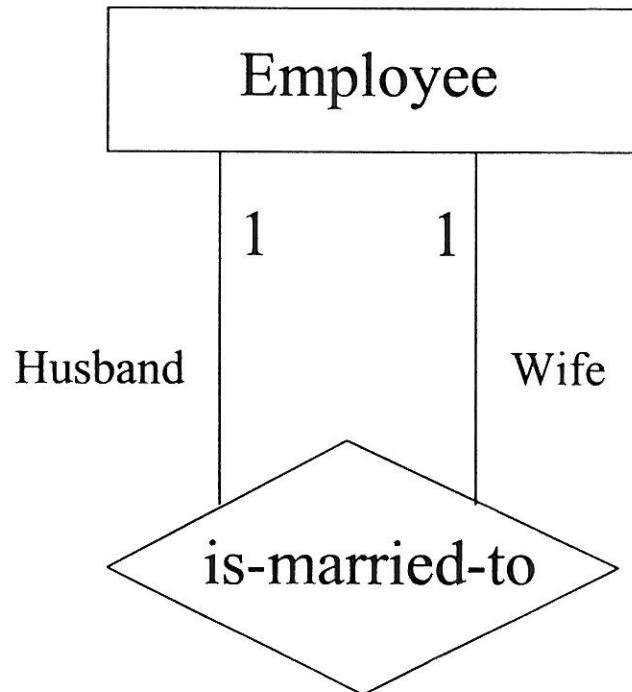
relationship set **embedded** in entity table  
corresponding to “child” and primary key of  
parent table appears as foreign key of child  
table



```
create table department
(dept_no integer,
dept_name char(20),
primary key (dept_no));
```

```
create table employee
(emp_id char(10),
emp_name char(20),
dept_no integer not null,
primary key (emp_id),
foreign key (dpt_no) reference
department
on delete set default on update
cascade);
```

# Recursive Binary Relationship



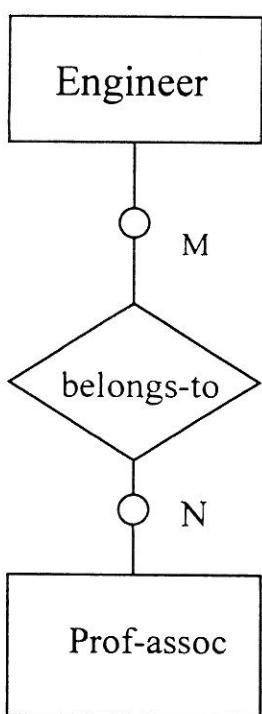
Create table **employee**

```
(emp_id char(10),  
emp_name char(20),  
spouse_id char(10),  
primary key (emp_id),  
foreign key (spouse_id) references  
employee on delete set null on  
update cascade);
```

What happened to Roles?

## 2.3 Many-to-many

relationship set represented as a **separate** table  
and the primary key on this table consists of the  
primary keys of the related entity tables



```
create table engineer
```

```
(emp_id char(10),  
primary key (emp_id));
```

```
create table prof_assoc
```

```
(assoc_name varchar(256),  
primary key (assoc_name));
```

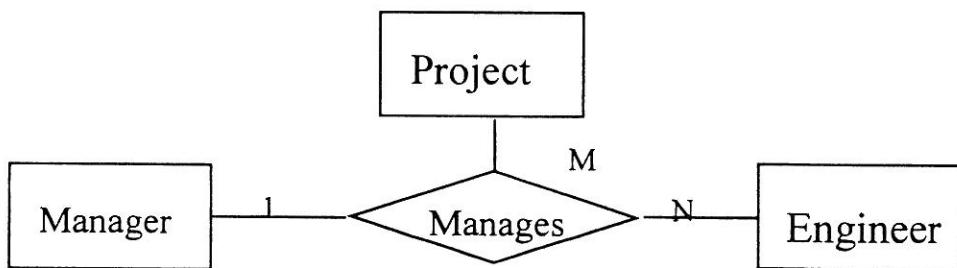
```
create table belongs_to
```

```
(emp_id char(10),  
assoc_name varchar(256),  
primary key (emp_id, assoc_name),  
foreign key (emp_id) references engineer  
on delete cascade on update cascade,  
foreign key (assoc_name) references  
prof_assoc  
on delete cascade on update cascade);
```

### 3. Ternary and N-ary relationships rule:

In all cases create a **separate** table containing the primary keys of all entities

However the primary key of this table may consist of less than n attributes



```
create table project
    (project_name char(20),
     primary key (project_name));
create table manager
    (mgr_id char(10),
     primary key (mgr_id));
create table engineer
    (emp_id char(10),
     primary key (emp_id));
create table manages
    (project_name char(20),
     mgr_id char(10) not null,
     emp_id char(10),
     primary key (project_name, emp_id),
     foreign key (project_name) references project
         on delete cascade on update cascade,
     foreign key (mgr_id) references manager
         on delete cascade on update cascade,
     foreign key (emp_id) references engineer
         on delete cascade on update cascade);
```

## 4. Generalization rule

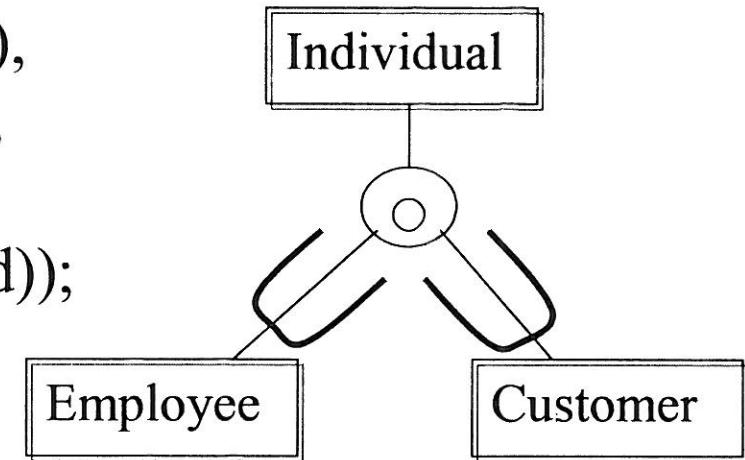
- create a separate table for the generic type and each of its subtype
- The supertype table contains the supertype key and all common attributes
- Each subtype table contains the supertype key [as primary and foreign key] and only the attributes that are specific to that subtype

# Generalization Abstraction Transformation Rules

```
create table individual (indiv_id char(10),  
                         indiv_name char(20),  
                         indiv_addr char(20),  
                         primary key (indiv_id));
```

```
create table employee (emp_id char(10),  
                      job_title char(15),  
                      primary key (emp_id),  
                      foreign key (emp_id) references individual  
                        on delete cascade on update cascade);
```

```
create table customer (cust_no char(10),  
                      cust_credit char(12),  
                      primary key (cust_no),  
                      foreign key (cust_no) references individual  
                        on delete cascade on update cascade);
```



## Create table

- Basic data definition statement in SQL'92

## Constraints on attributes

- Not null
- Unique for candidate keys
- Primary key => implies not null & unique
- Foreign key

## How to enforce the foreign key rule ?

- two options :

(A) Reject any operations that result in **illegal state**

or

(B) System performs additional **compensating actions** to guarantee that overall result is a **legal state**

# Default Policy: Rejecting Violating Modifications

- 1) Try to insert a new **employee** tuple whose *dept\_no* is not known or is NULL.
- 2) Try to update an **employee** tuple to change the *dept\_no* to a non-NULL value that is not a legal value in **department** table.
- 3) Try to delete a **department** tuple and its *dept\_no* appears as a component of an **employee** tuple(s).
- 4) Try to update a **department** tuple in such a way that we change the *dept\_no* and the old *dept\_no* appears as a component of an **employee** tuple(s).

**ALL ARE REJECTED**

## SQL '92 solution

- On update

**Cascade -** The update operation on the PK in the referenced table cascades to all tuples in the referencing table with matching FKS.

**Set-Null -** Foreign keys are set to NULL when they match the old PK of the updated tuple in the referenced table.

**Set-Default -** Foreign keys are set to default when they match the old PK ...

[Note: broader interpretation of NULL]

- On Delete

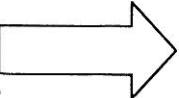
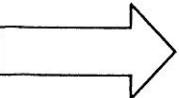
**Cascade -** The delete operation on the referenced table cascades to all tuples in the referencing table with matching FKS

=> One user operation may trigger many other operations

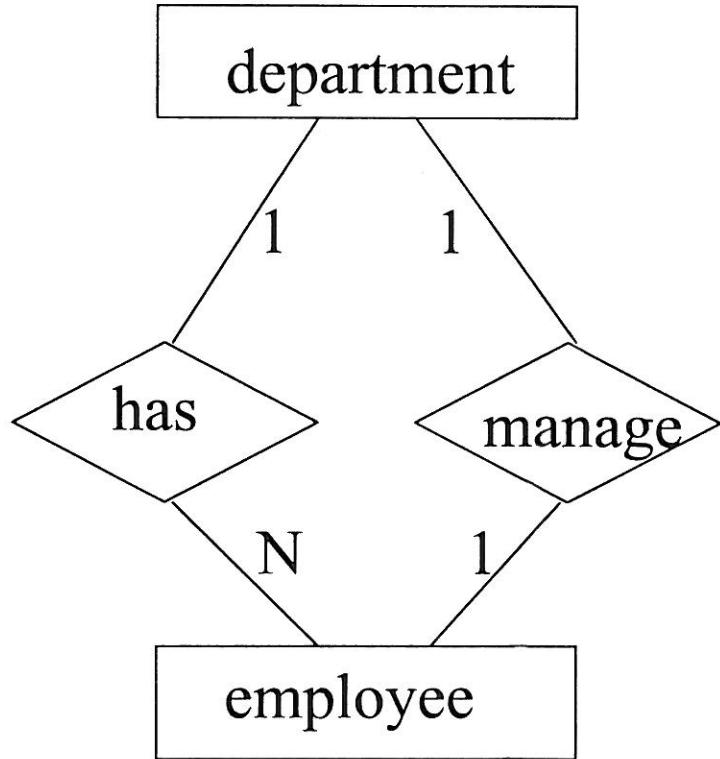
**SET-NULL ...**

**Set-Default ...**

## Previous Solution Handles only Modifications of Types 3 and 4

- 1) *insert into employee  
values ('chair01', 'abejones');*  System cannot use NULL  
for foreign key
  
- 2) *update employee  
set dept\_no = 14  
where emp\_id = 'bigboss';*  If dept\_no = 14 is a new  
value not in **department**  
table action is rejected

# Embed Two One-to-many Relationship



```
create table department  
(dept_no integer,  
dept_name char (20),  
emp_id char (10),  
primary key (dept_no),  
foreign key (emp_id) references employee  
on delete set default on update cascade);
```

**Question:** Can we insert tuples  
with new *dept\_no* or new *emp\_id*?