

EECS 495 – Introduction to Database Systems

Lecture – 2

Prepared by Mas-ud Hussain

TA Office Hours

- Will start from **9/25 (Friday)**

Tuesday: 5:00 – 6:00 pm

Wednesday: 4:00 – 5:00 pm

Friday: 2:30 – 3:30 pm

(and/or, *appointment via email*)

Location: Tech L580

Email: mmas-ud.hussain@u.northwestern.edu
mmasud.hussain@gmail.com

- There will be additional **grader(s)** to help in grading.

Last Week...

- Introduction
- File Systems vs. DBMS
- Core Database Functionalities
 - Data Independence
 - Declarative Querying
 - Transactions (contd..)
- Metadata (contd..)

Transactions: *An Execution of a DB Program*

- Key concept is transaction, which is an *atomic* sequence of database actions.
 - Insert
 - Delete
 - Update
 - Reads/Writes
- Each transaction, executed completely, must leave the DB in a consistent state if DB is consistent when the transaction begins.

Transactions: *Example of Inconsistency*

Instructors

ID	Name	Designation
1	Peter Scheuermann	Professor
2	Douglas Downey	Associate Professor
.....

DELETE
ID 1

Courses

ID	Name	Instructor ID
EECS 495	Intro to DB	1
EECS 317	Data Management	2
EECS 230	Intro to Programming	3
.....

Inconsistency

Transactions: *Example of Inconsistency*

Account

ID	Type	Balance	Credit Line	Total Cashback
1	Gold	600	2000	21
2	Premium	220	4000	12
.....

Transaction

ID	Account ID	Amount	Cashback (5%)
1	2	50	2.5
2	1	40	2
.....

 INSERT new transaction

3	1	100	6
---	---	-----	---

Transactions: *Maintaining Consistency*

- Users can specify some simple *integrity constraints* on the data, and the DBMS will enforce these constraints. (1st example)
- Beyond this, the DBMS does not really understand the *semantics* of the data. (e.g., it does not understand how the interest on a bank account is computed). (2nd example)
- Thus, ensuring that a transaction (run alone) preserves consistency is ultimately the user's responsibility!

Queries in a DBMS

- DBMS provides a specialized language for accessing data:
 - Query Language
- Query language can be further classified into:
 - DML - Data Manipulation Language (Insert/Delete/Update)
 - DDL - Data Definition Language (Create Tables)
 - DCL - Data Control Language (Procedural Tasks)

Queries in a DBMS

- Standard for relational DBMS: **SQL**
 - Based on formal query languages: **Relational Algebra and Relational Calculus**
- Queries are evaluated as efficient as possible
 - Huge influence of physical design

Queries: “What”, not “How”

- It is convenient to indicate declaratively *what* information is needed, and leave it to the system to work out *how* to process through the data to extract what you.
 - Just like *programming*, e.g., if..else, for, while, printf, etc. keywords to declare what to do.
- A query language is based on declarative logic.
 - SELECT name FROM customer WHERE id=1 (*SQL*)

DB System Requirements



- **High Availability:** on-line => must be operational while enterprise is functioning (Page not loading!)
- **High Reliability:** correctly tracks state, does not lose data, controlled concurrency (Lose a Facebook photo)
- **High Throughput:** many users => many transactions/sec (1.5+ billion total users)
- **Low Response Time:** on-line => users are waiting (Respond within ms)
- **Security:** sensitive information must be carefully protected since system is accessible to many users (passwords, security question, privacy controlled images)

Roles in DB Management

- **System Analysts**
 - Specifies system using input from customer; provides complete description of functionality from customer's and user's point of view
 - Conceptual database design
- **Database Designer**
 - Specifies structure of data that will be stored in database (logical & physical database schemas)
- **DB Application Programmer**
 - Implements application programs (transactions) that access data and support enterprise rules
- **DB Admin**
 - Maintains database once system is operational
- **End-Users**
 - Often unaware that they are dealing with data in a DBMS

Summary of Introduction

- DBMS used to maintain & query large datasets that are shared by many application programs/users
- Some powerful ideas:
 - Program-Data Independence
 - Controlled Data Redundancy
 - Declarative Queries
 - Transactions
- Every ‘knowledge worker’ or scientists needs database know-how, as do all IT experts (application developers, software engineers, system analysts, ...) - not just DBAs
- Databases are one of the broadest and most useful areas in CS and IS

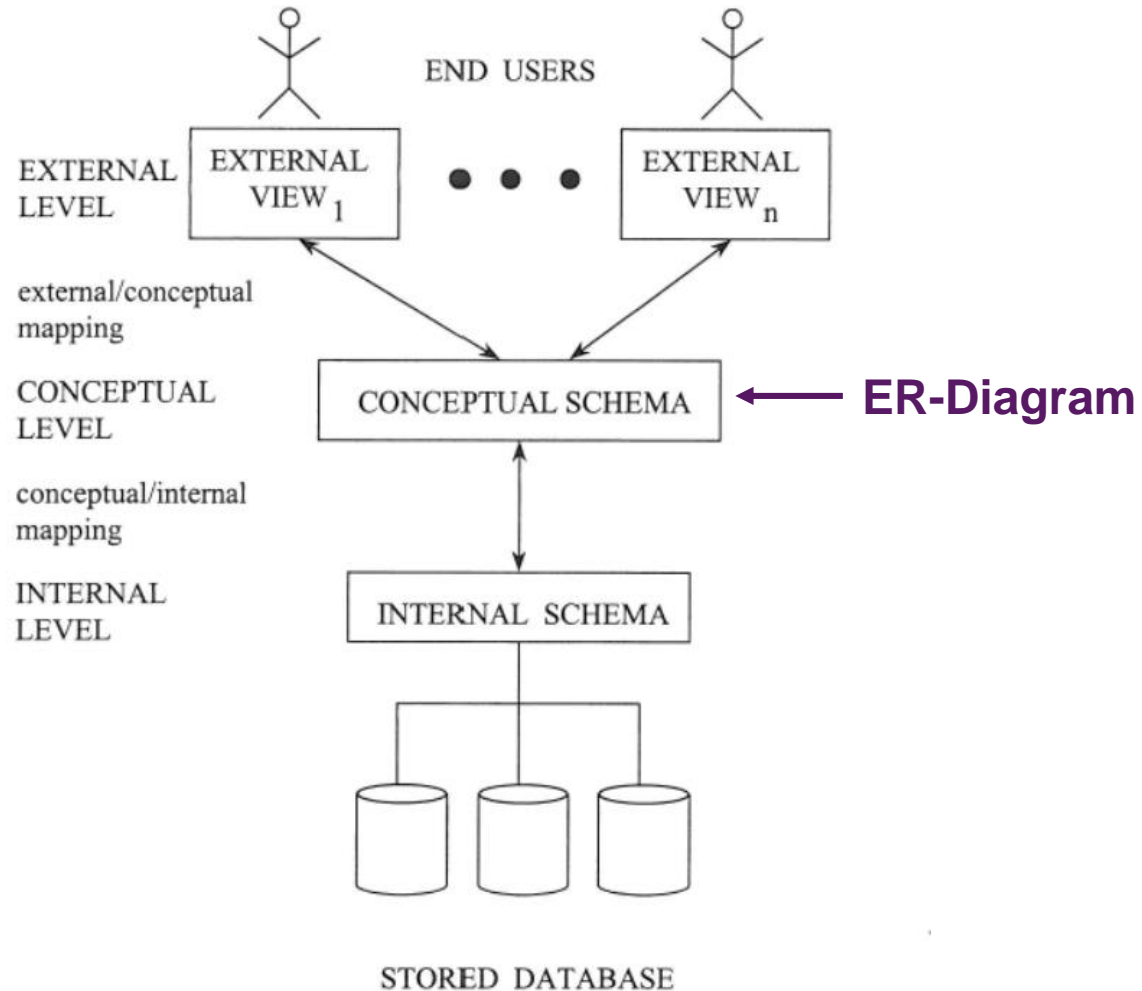
ENTITY RELATIONSHIP DIAGRAM (ERD)

Data Modeling

Outline

- Introduction --- Database Architecture Framework
- **Data Modeling with Entity Relationship Diagram**
- Introduction to the Relational Model
- Transformation of ER-Schema to Relational-Schema
- Normalization in the Relational Model
- Overview of SQL

ANSI/SPARC Architecture



Model

- A formal representation of user specification/requirements for a project.
- Description of the real world!
- An important tool in the communication among users, system analysts, database administrators.

Entity Relationship Model

- In this class, we will confine to the **ER = Entity-Relationship model**, where the basic “players” are:
 - Entities (Entity-Set)
 - Their Attributes
 - Relationship among the entities
 - The attributes of the relationships
- There are other modeling tools on the market, with a similar nature/objectives, but used in different context(s):
 - **UML (Unified Modeling Language)** used for different software purposes

Entity - Definition

- **Entity:** Anything that can be distinctly identified, i.e., has existence in the real world. E.g., a person (Barack Obama).
- **Entity-Set:** a group of entities with identical properties. E.g., Politician.
- Barack Obama (the entity) is a member of the entity-set Politician.
- Sometimes the word Entity is used to refer to Entity-Set.

Entity - Example

- **Database of a University (e.g., Northwestern):**
 - Student
 - Faculty
 - Non-Academic Staff
 - Department
 - Courses, etc.
- **Database of a Bank (e.g., Chase):**
 - Customers
 - Employee
 - Manager
 - Branches
 - Account
 - Deposit/Transaction, etc.

Entity - Representation

- An entity is represented in the diagram as below:



Employee

Student

Attributes - Description of Entities

- **NOTE:** entity classes need not be “fully disjoint” (in the sense of not having anything in common whatsoever...):
 - A very same person can be **both employee and customer** of the bank.
- Essentially, one can think that every **entity** is represented by a set of **Attributes** (table columns):
 - (descriptive) set of **properties** that every instance of that entity class **should** have

Attributes - Example

- An employee can have:
 - Name
 - Age
 - Salary
 - Join Date
 - Address
 - Phone No., etc.
- A student can have:
 - Name
 - Net-id
 - Major
 - CGPA
 - Degree Level, etc.

Attributes - Example

- We can represent an entity-set as follows:

Name_Of_Entity (Attribute 1, Attribute 2,Attribute N)

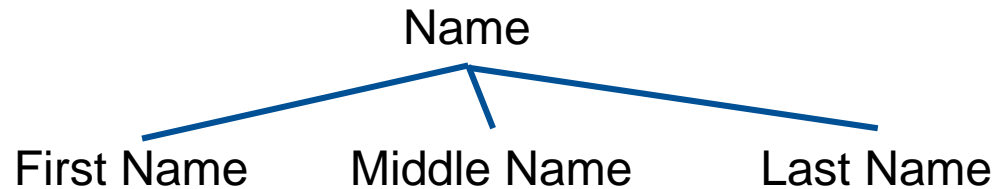
- **Example 1:** Employee (Name, Age, Salary, Join Date, Address)
- **Example 2:** Student (Net-id, Name, Major, CGPA, Degree-Level)

Attributes – Types (Single-valued vs Multi-valued)

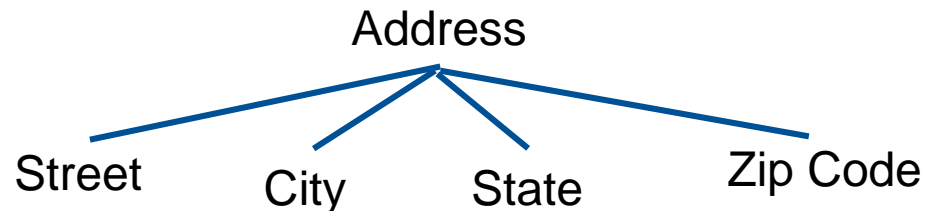
- Most attributes are Single-valued, e.g., **Net-id, Salary, etc.**
- The nature of certain attributes is such that they have > 1 from the domain (multi-valued).
- **Example1:** Phone Number --- A person may have multiple phone numbers.
- **Example2:** Degrees --- A person can have multiple degrees (B.Sc., M.Sc., Ph.D., etc.)

Attributes – Types (Simple vs Composite)

- **Composite attributes** can always be split into a collection of attributes that are “atomic” (i.e., their domain is “simpler”)
- **Example1:**



- **Example2:**




Attributes - Domain

- For each individual attribute, there has to be a specification of the (set of the) “permissible” values, called domain. In addition, for each domain, we need:
 - **Type** of the domain
 - **Restrictions** on the **values** from the respective domain
- **Example1:** Attribute :- **Name**, Domain :- **String**, Restriction:- **Maximum 30 characters**.
- **Example2:** Attribute :- **Age**, Domain :- **Number**, Restriction:- **[16, 70]**.

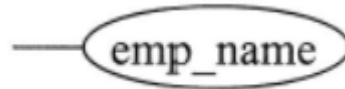
Attributes – NULL Values

- **NULL** is a special **symbol** that denotes the situation in which a particular instance does not have a value for the given attribute. It can indicate:
 - Non-Applicable values
 - Single family home has no apt no
 - Missing Data
 - Not Known
 - We don't know the state where Mary smith lives.

Attributes – Representation

- Attributes are represented as below (oval-shaped) — 

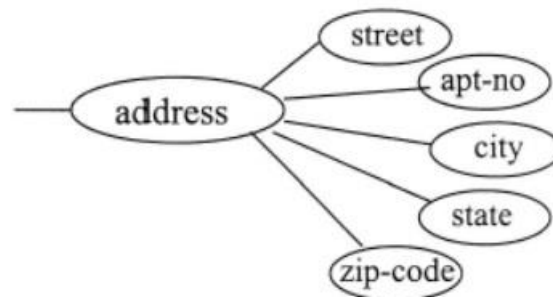
- Example:



- Multi-valued:



- Composite:

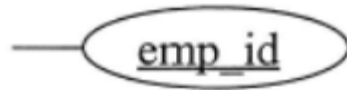


Attributes – Primary Key

- How to retrieve a particular entity from an entity-set? (i.e., retrieve a tuple from a table)
- **Primary Key:** An attribute (or, a set of attributes) that uniquely identifies an entity within an entity-set.
- **Example 1:** In case of the **student** entity, **net-id** must be unique for all the students. Thus, **net-id can be a primary key** of the student entity.

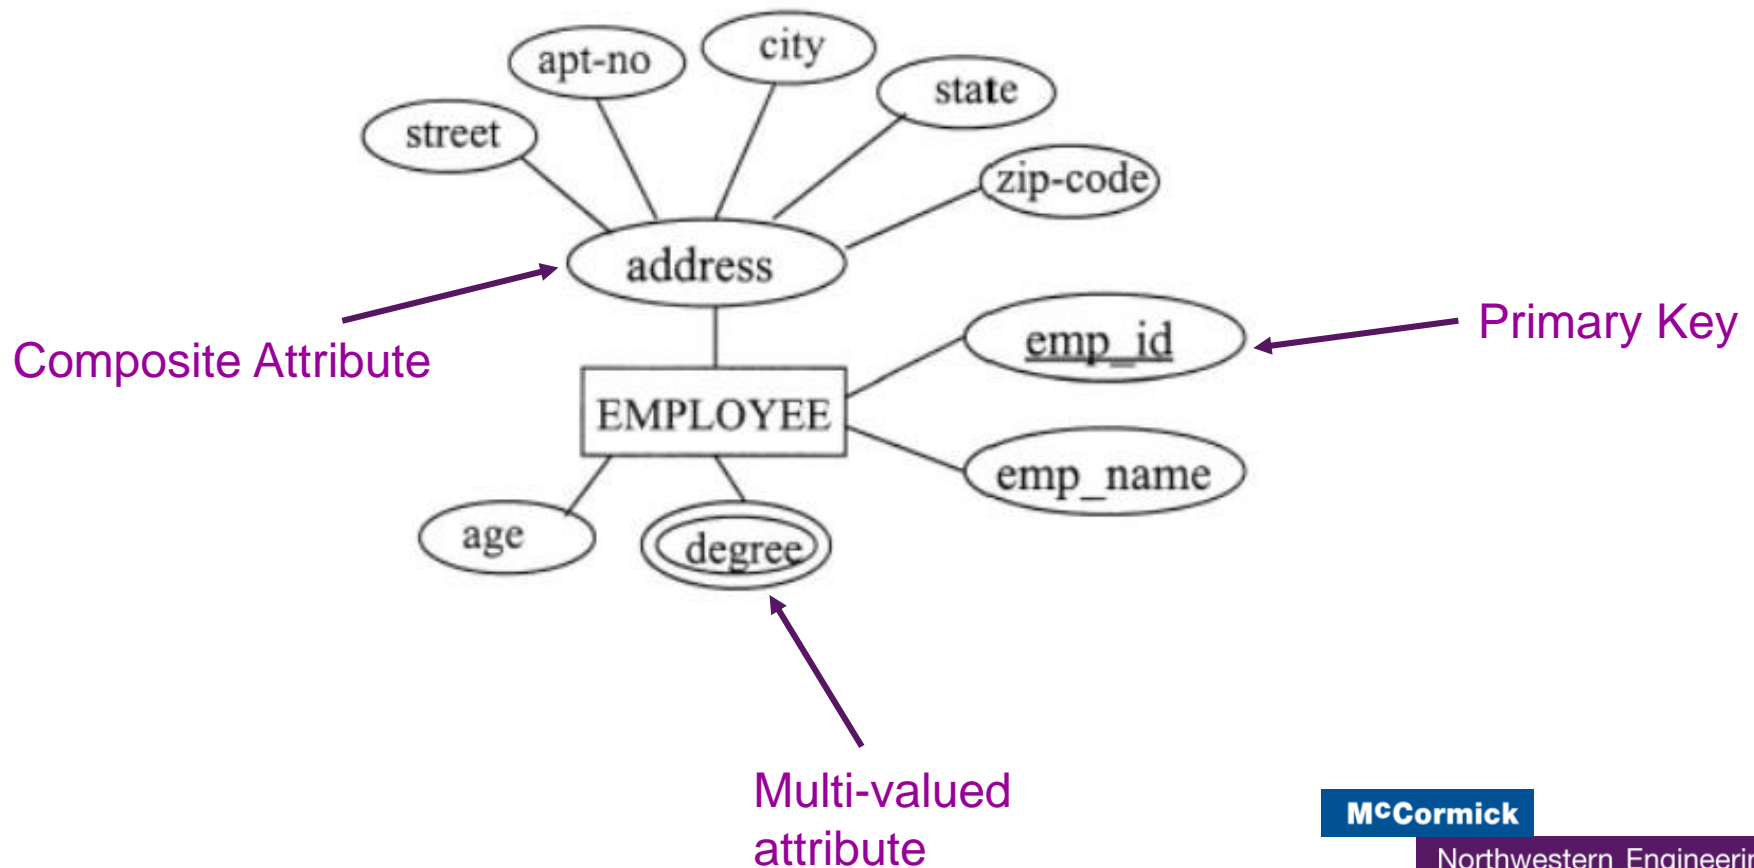
Attributes – Primary Key

- **Example 2:** In case of the **course** entity, the combination of two fields: **department** and **course number** must be unique for all the courses. E.g., “EECS 339”. Thus, (department, course number) can be a primary key for the **course** entity.
- Each entity(table) must have a primary key.
- Primary key is represented in the model as below:



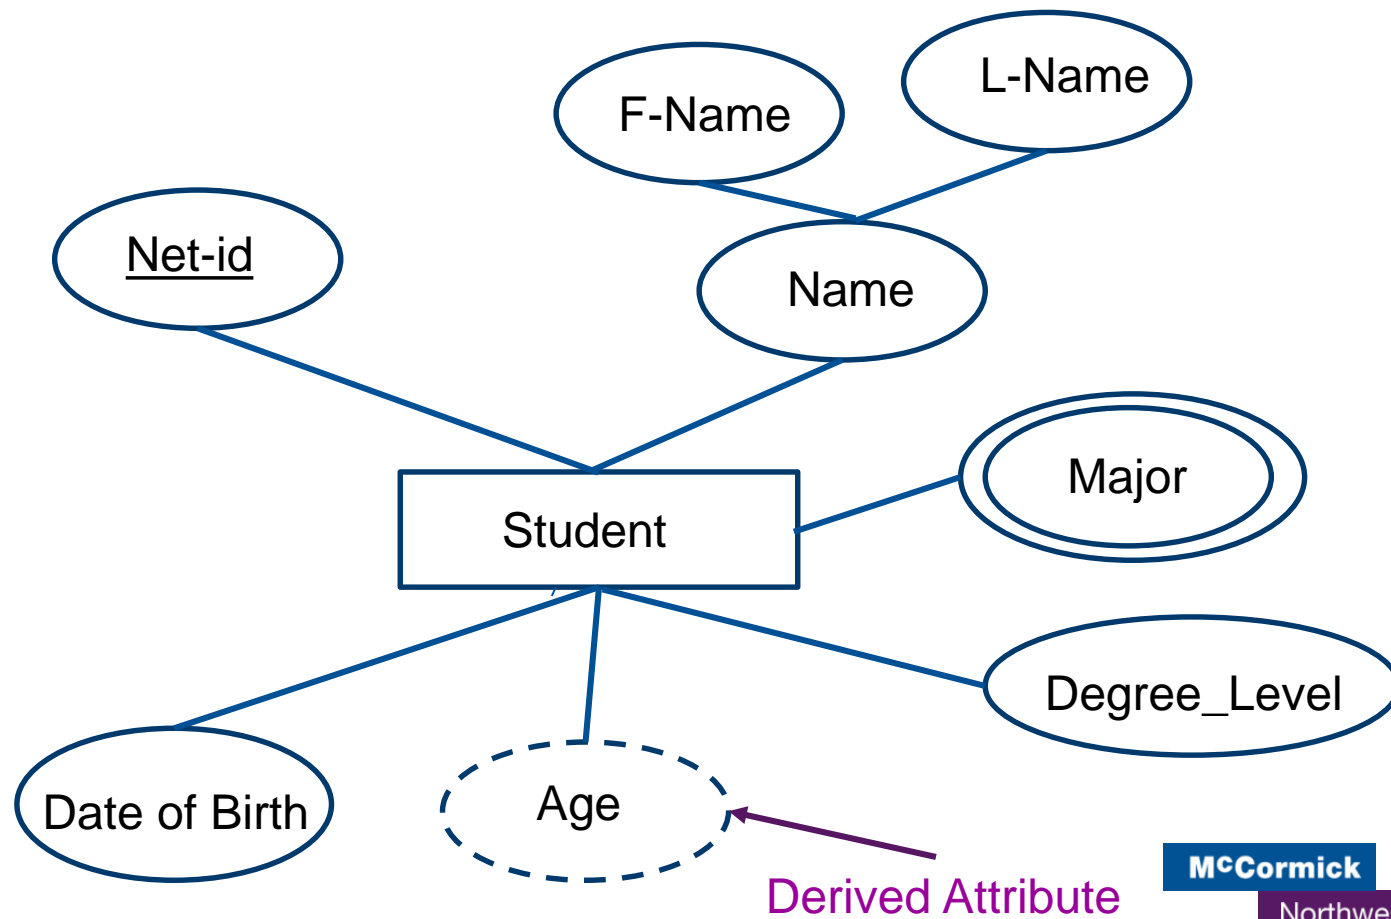
Example Representation

- Employee (emp-id, emp-name, address, degree, age)



Example Representation

- Student (net-id, first_name, last_name, major, degree_level, date_of_birth, age)



Modeling Relationship

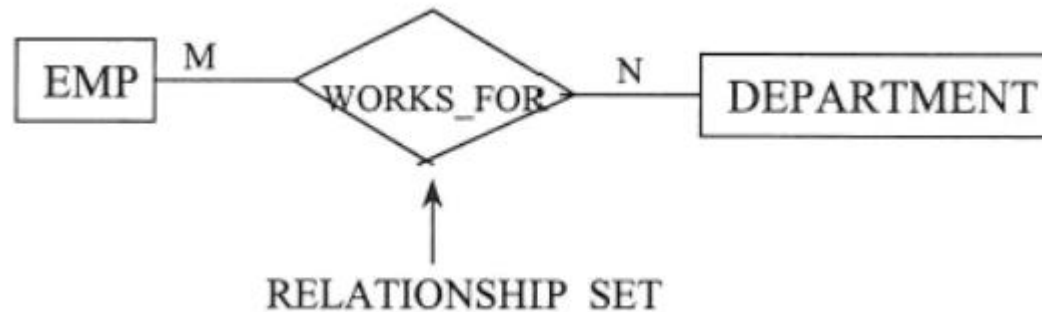
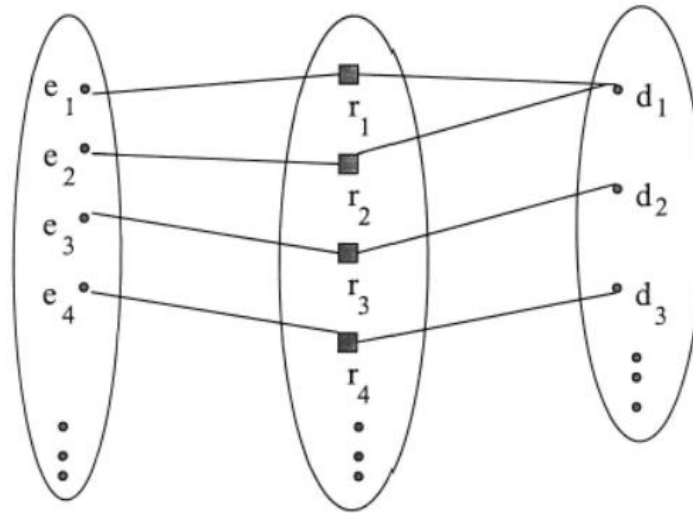
- **Relationship:** A (possible) association among the instances of two (or more) entities.
- **Example:**

Hayes
customer entity

depositor
relationship set

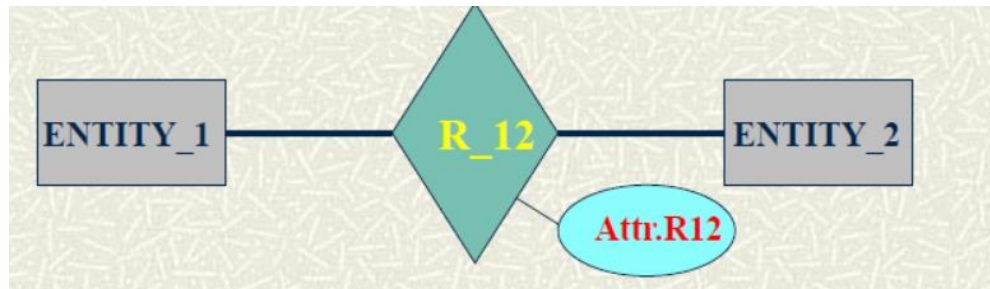
A-102
account entity

Modeling Relationship - Examples



Modeling Relationship

- **Relationship (Type):** A description of the possibility of two (binary) or more (ternary, quaternary,...) entity classes having some association among their instances. **NOTE: may have attributes of its own.**
- **Relationship Instance:** An association between two (or more) entities. Note that in an actual database, these will be rows/tuples of the tables.
- Symbol:



Modeling Relationship

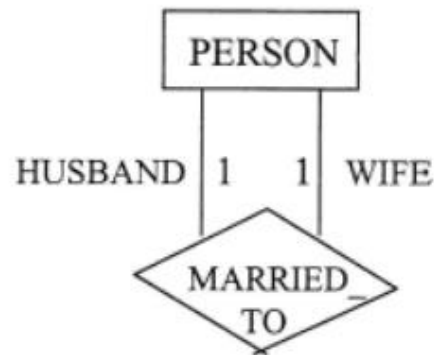
- Relationships are characterized in terms of:
 - *Degree*
 - *Cardinality*
 - *Constraints*
 - *Attributes*

Degree of Relationship

- **Degree**: No. of participating entity-sets.
- Binary Relation:

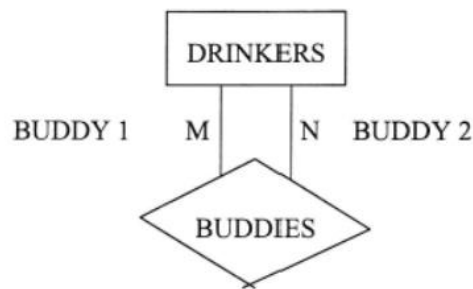


- Recursive Binary (Relation between same entity-sets):



Degree of Relationship

- Recursive Binary (Relation between same entity-sets):

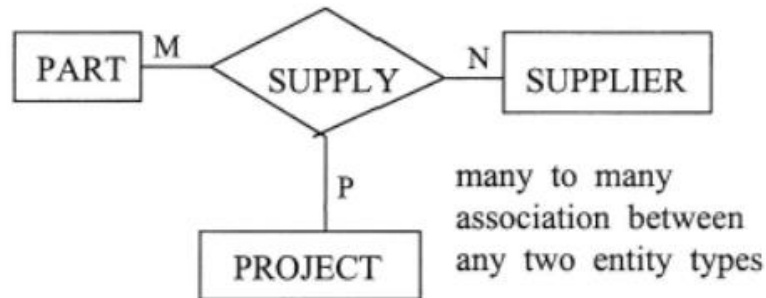


Buddy 1	Buddy 2
d ₁	d ₂
d ₁	d ₃
d ₂	d ₁
d ₂	d ₄

- Buddies is *symmetric*, married is not.
- No way to model *symmetry*.
- Design Question:** Should husband and wife be replaced by spouse?

Degree of Relationship

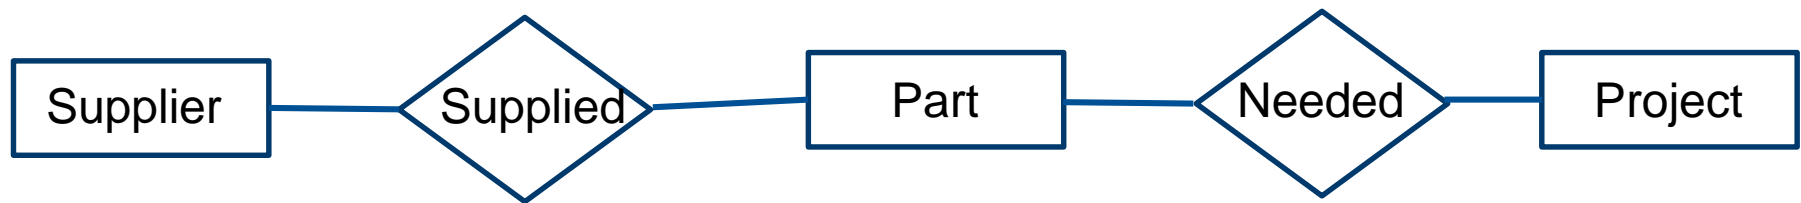
- Ternary Relationship:



PART#	SUPPLIER#	PROJ#
25	4	1
25	5	2
10	4	2
17	4	3
17	2	1
17	5	1

Degree of Relationship

- Ternary (and higher) Relationship should be avoided if possible (into binary relations):



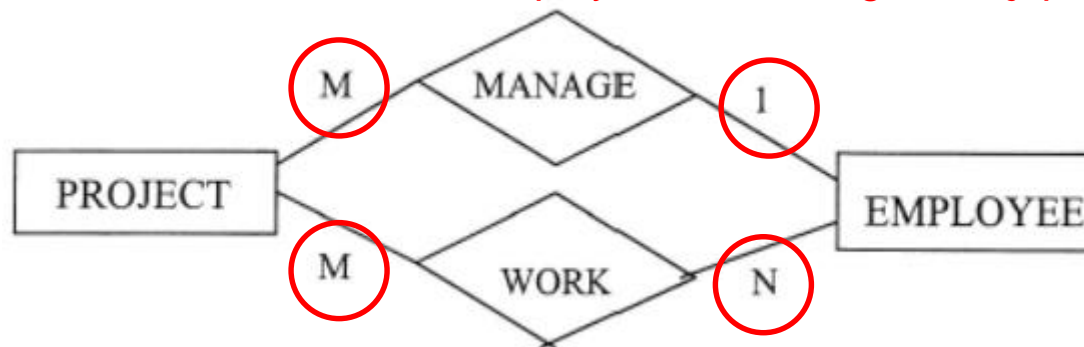
- But, it is unavoidable when two binary relationships are not equivalent
 - E.g., Supplier 4 supplied Part 25 for Project 1, but Supplier 5 supplied Part 25 for Project 2.

PART#	SUPPLIER#	PROJ#
25	4	1
25	5	2

Cardinality of Relationship

- **Cardinality Ratio**: Specifies the number of relationship instances that an entity can participate in.

Each project is managed by **one** employee (Manager),
But an employee can manage **many** projects

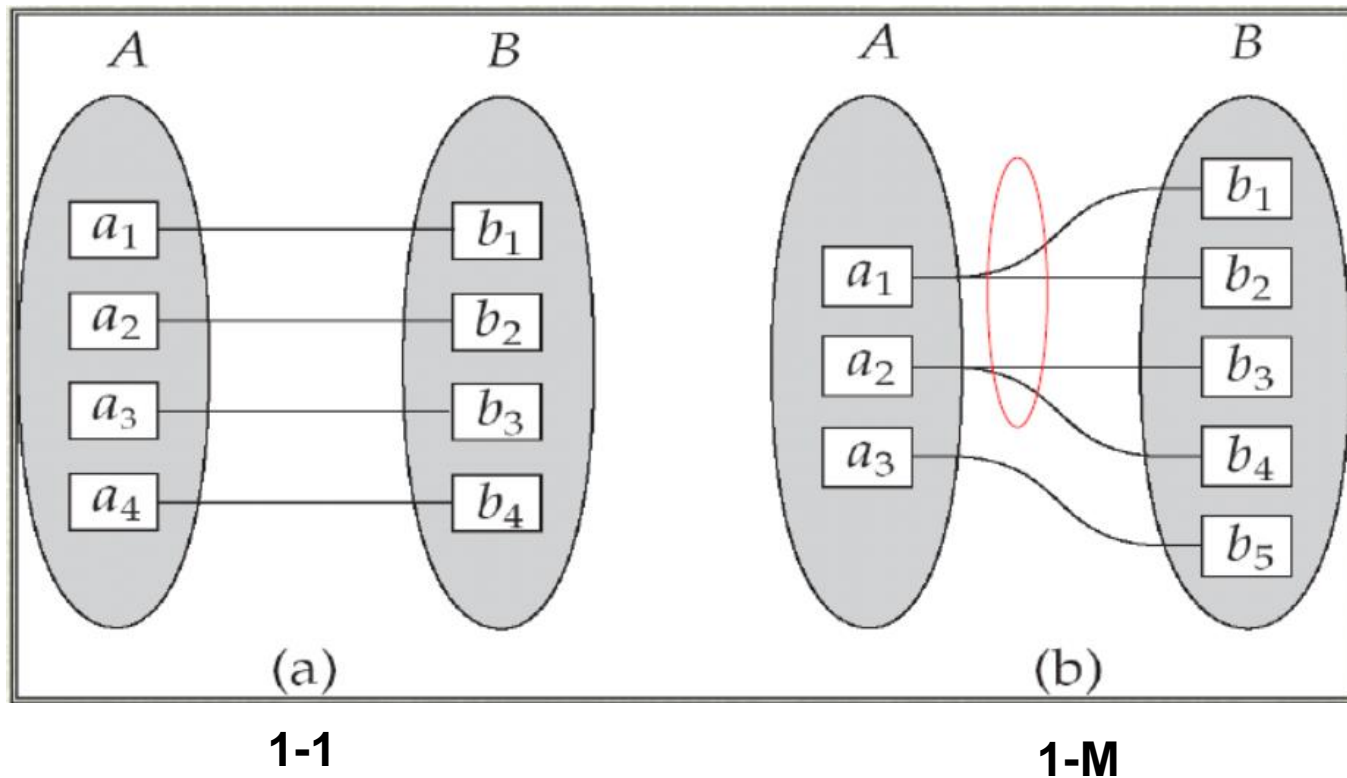


Each project has **many** employees,
And an employee can work on **many** projects

- Possible mappings are ---1:1, 1:M, M:1, M:M (or, M:N)

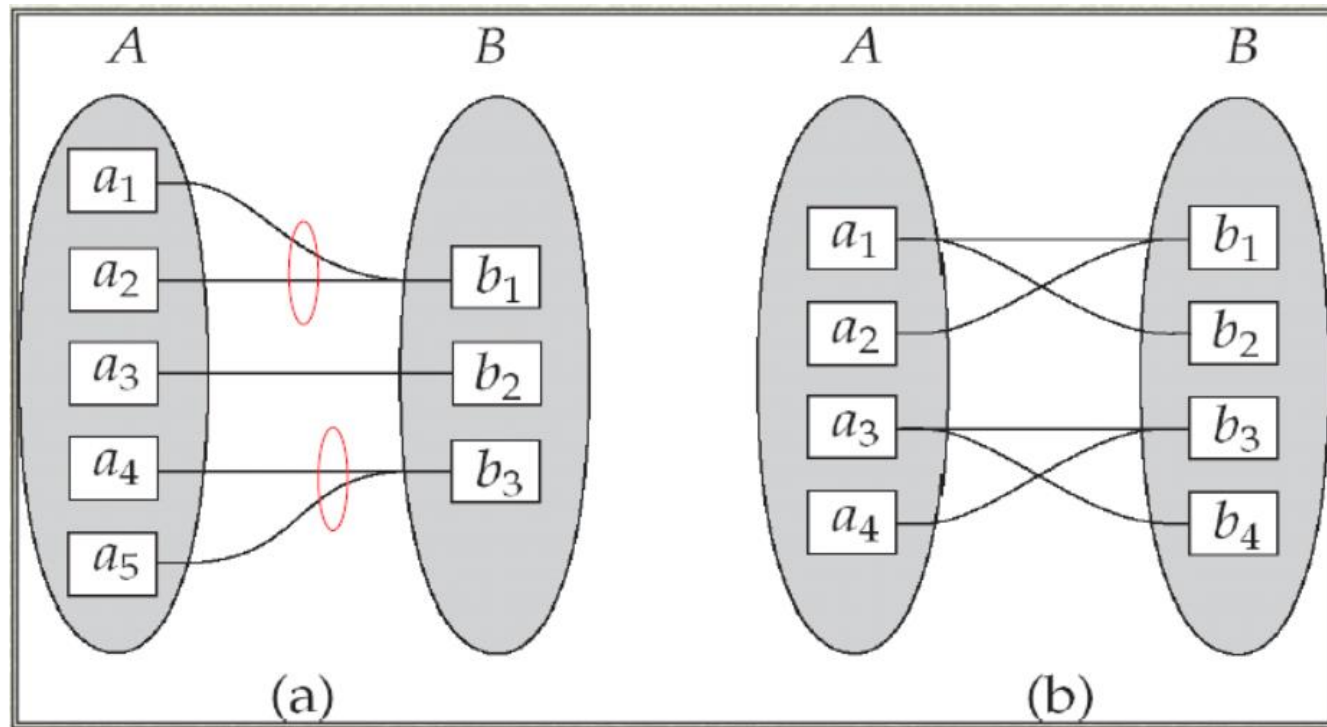
Cardinality of Relationship

- **Important:** *one* and *many* refers to the maximum cardinality.



Cardinality of Relationship

- **Important:** It could quite be the case that some elements in A and/or B **are not mapped** to any element(s) in the other set

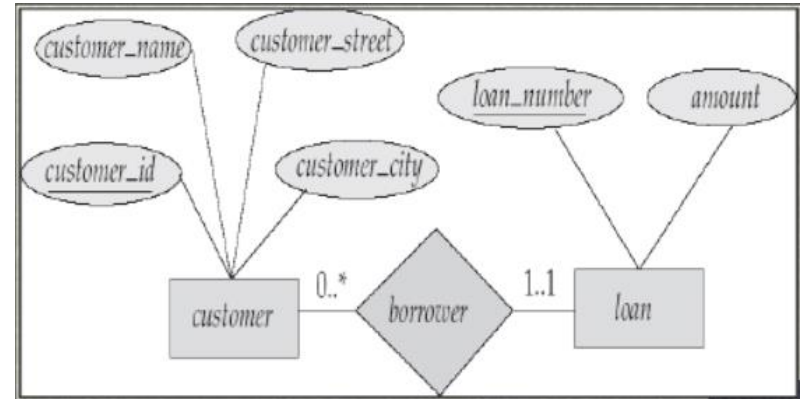
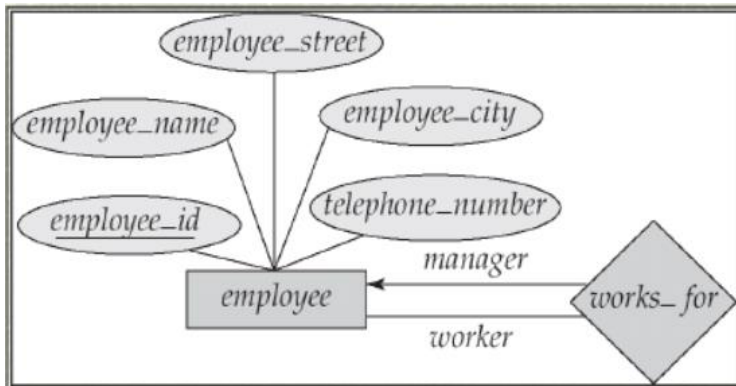


M-1

M-M

Cardinality of Relationship

- **Important:** The notation is not fully standardized, e.g., other notations are also used.

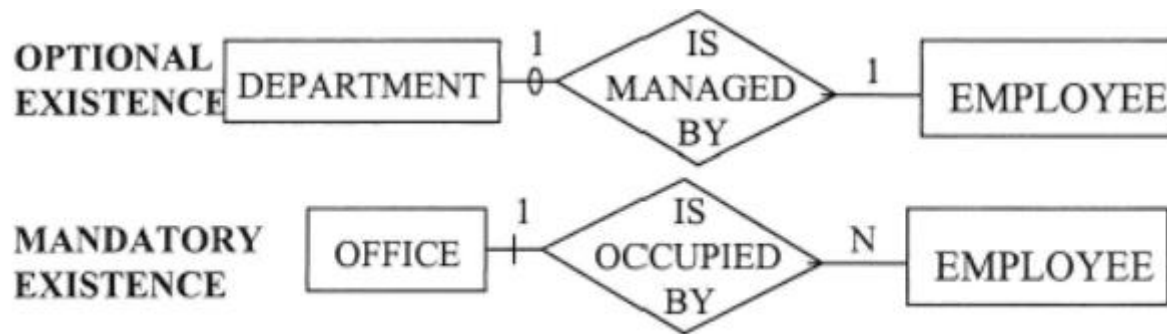


→ One

— Many

Constraints of Relationship

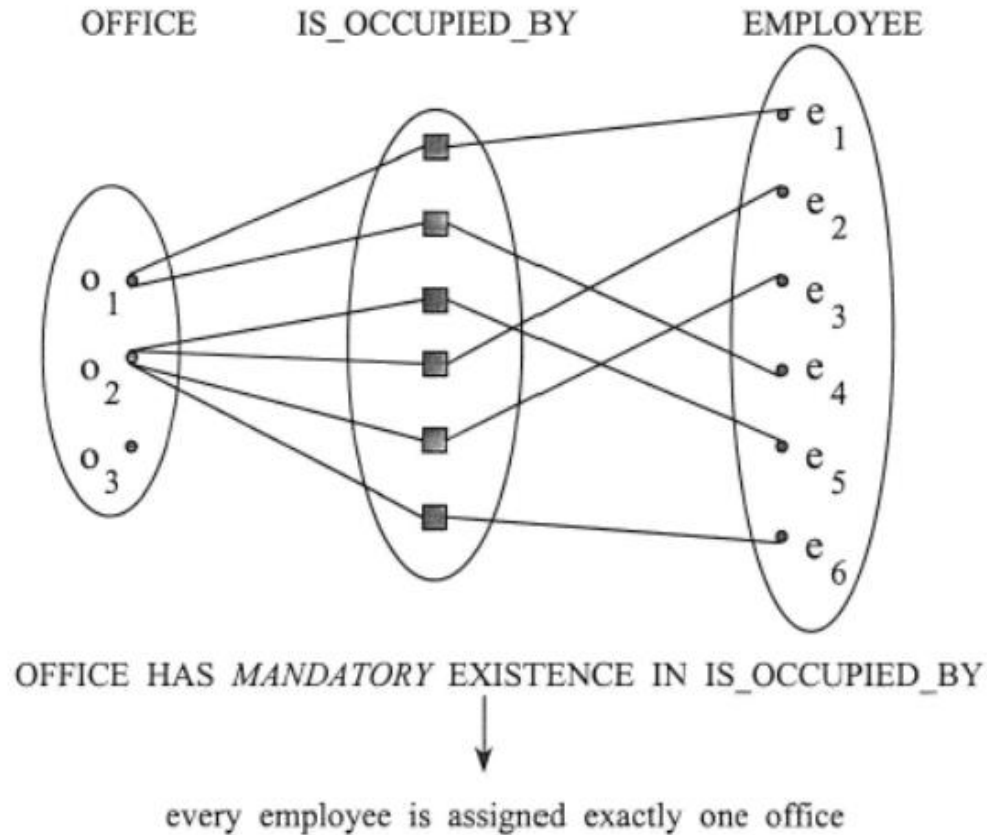
- MAY vs. MUST (AKA “partial” vs. “total” participation constraint)



OPTIONAL EXISTENCE: defines a minimum cardinality of zero (Partial)

MANDATORY EXISTENCE: defines a minimum cardinality of one (Total)

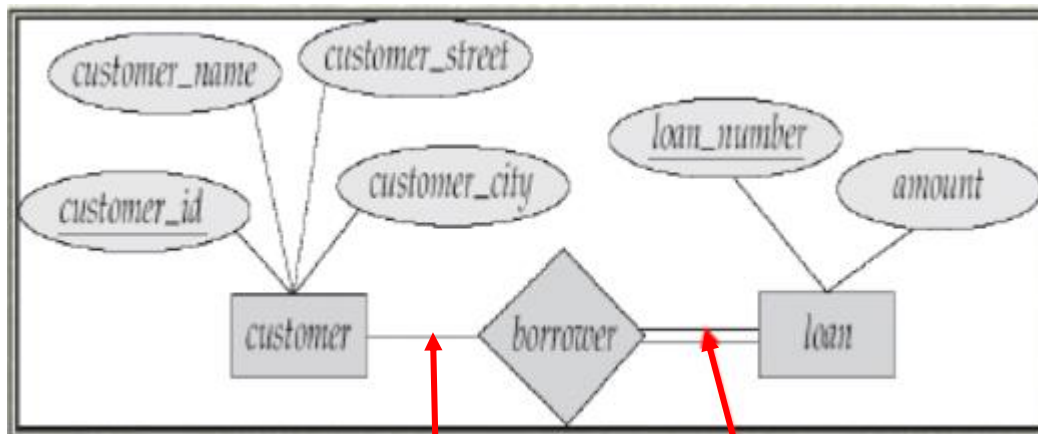
Constraints of Relationship



Same is not true for the Department and Employee in the previous example

Constraints of Relationship

- Participation of **Loan** is **total**, in the sense that every instance of **Loan must** be in a relationship with **some instance** from **Customer** (not vice versa, though)

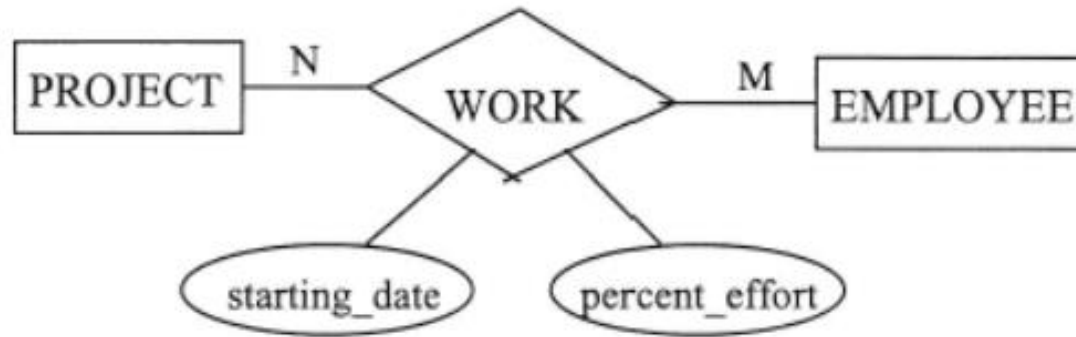


Partial/optional

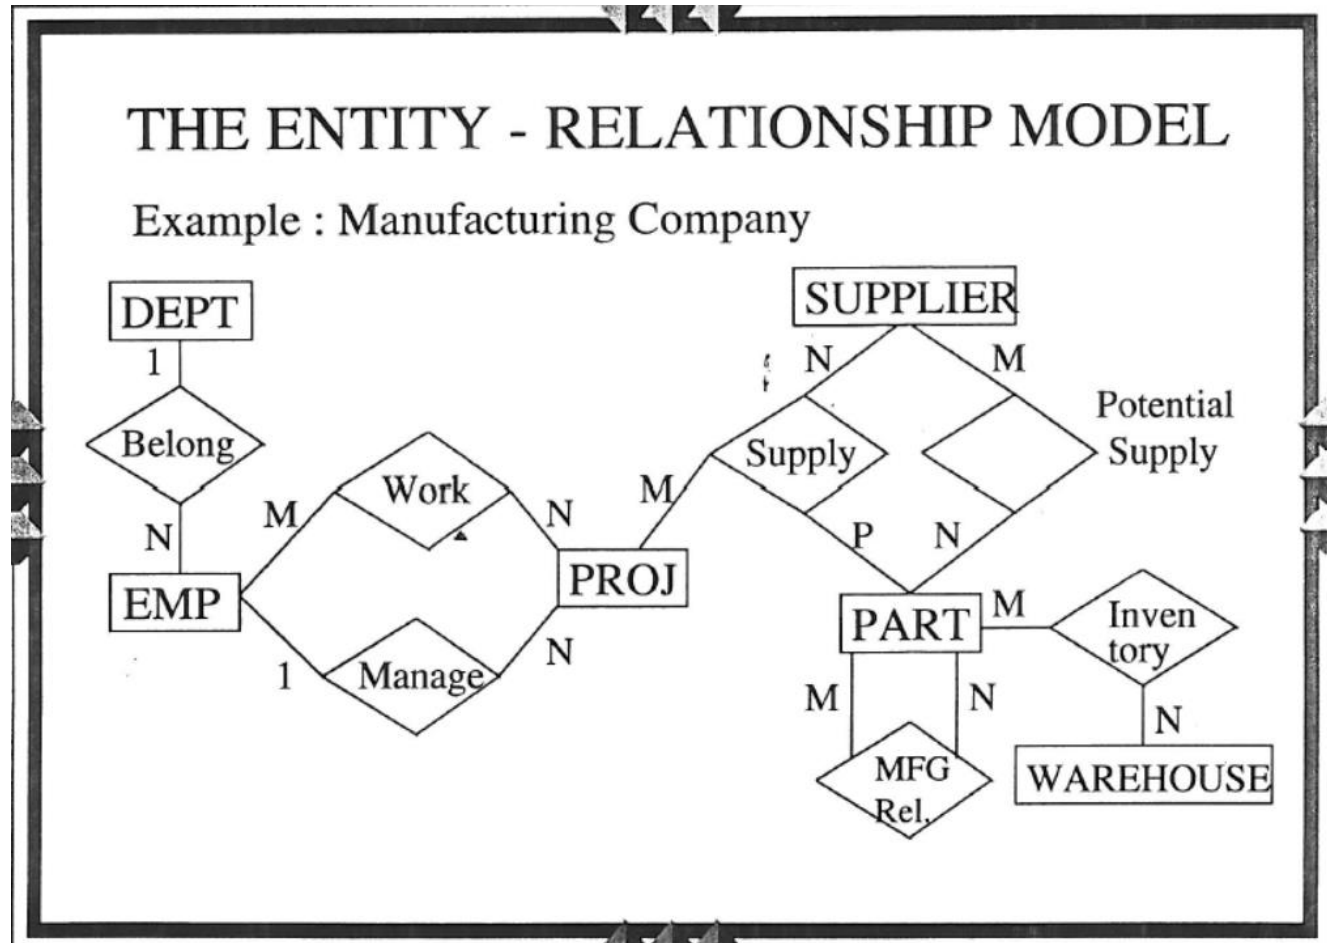
Total/mandatory

Attributes of Relationship

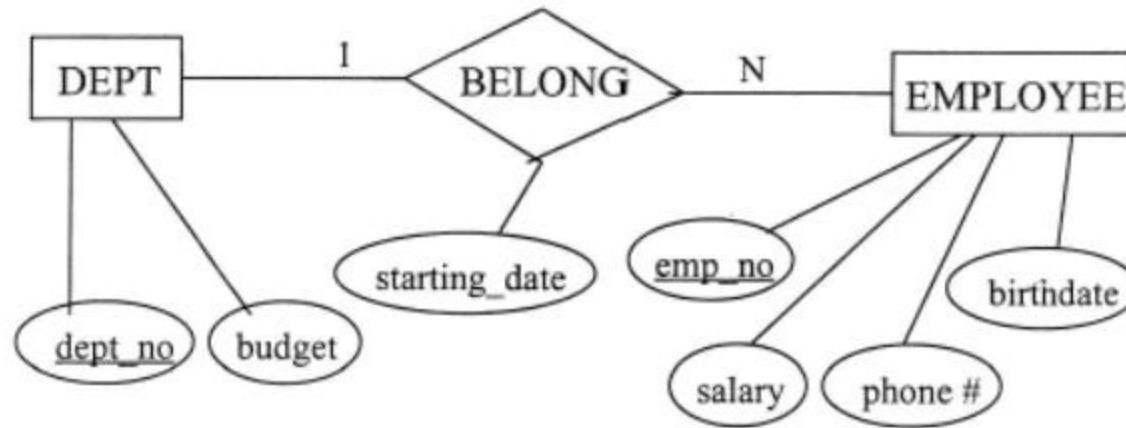
- Relationships can have **attributes** too!



ERD - Examples



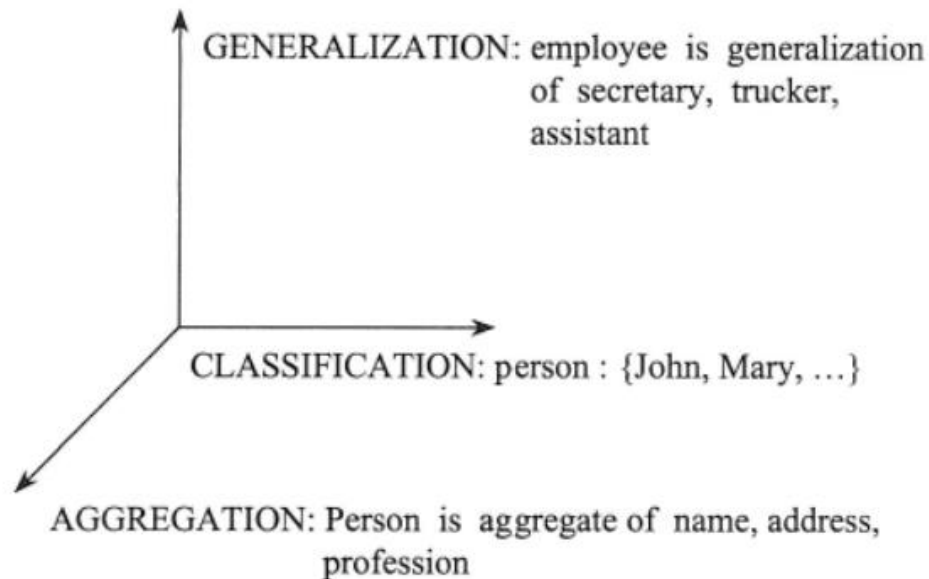
ERD - Examples



NOTE: attributes *birthdate* and *starting_date* are defined on the same value-set

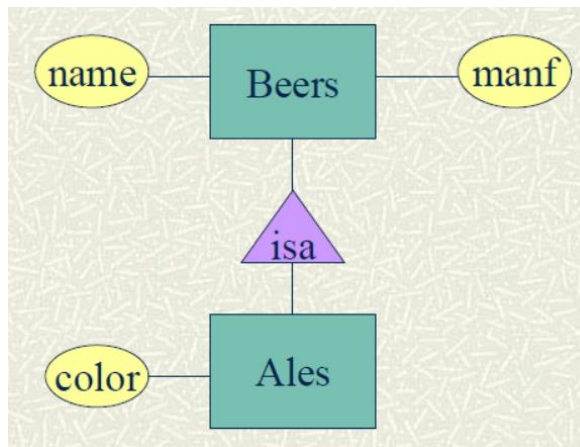
ERD – Advanced Topics

- **Classification:** Collection of objects with similar properties viewed as a class.
- **Aggregation:** Abstraction by which an object is constructed from its constituent objects.
- **Generalization:** Set of dissimilar objects are considered as a single generic object.



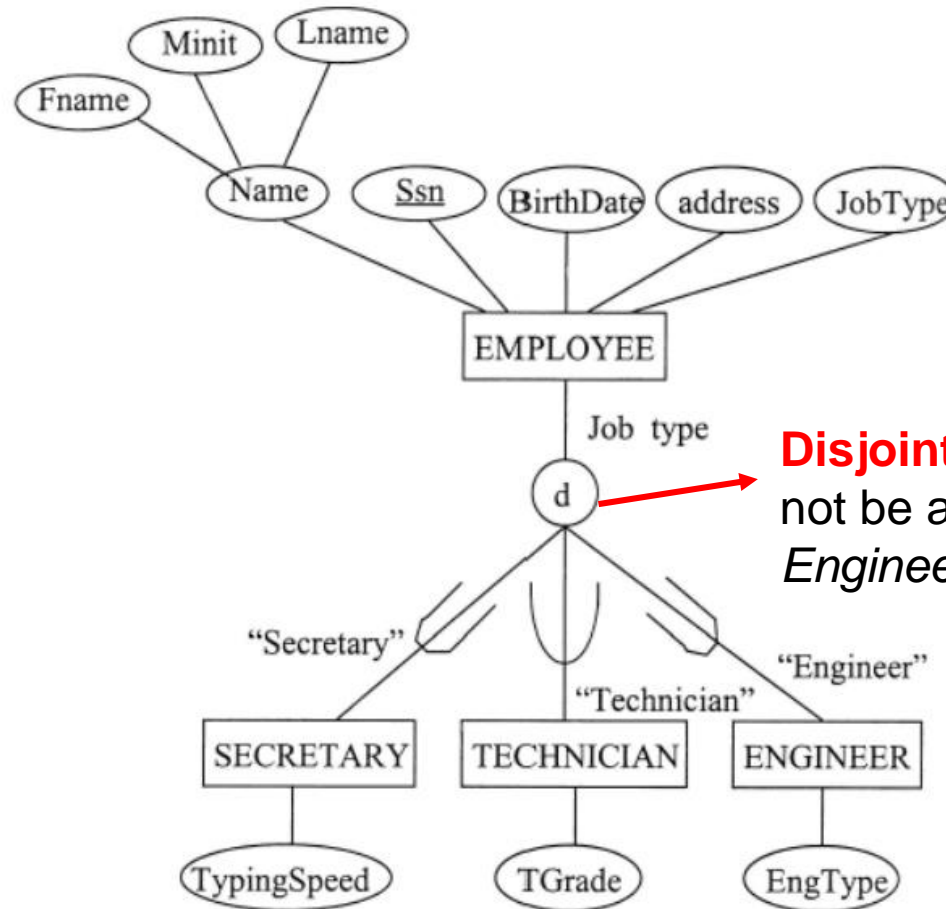
Generalization Hierarchy

- **Class/Subclass Hierarchy:** Similar to object-oriented programming concept *inheritance*, e.g., in Java, C++, etc.
- **Example:** Ales are a kind of beer.
 - Not every beer is an ale, but some are.
 - Let us suppose that in addition to all the *properties* (attributes and relationships) of beers, ales also have the attribute *color*.



NOTE: there may be other “conventions” to indicate subclasses...

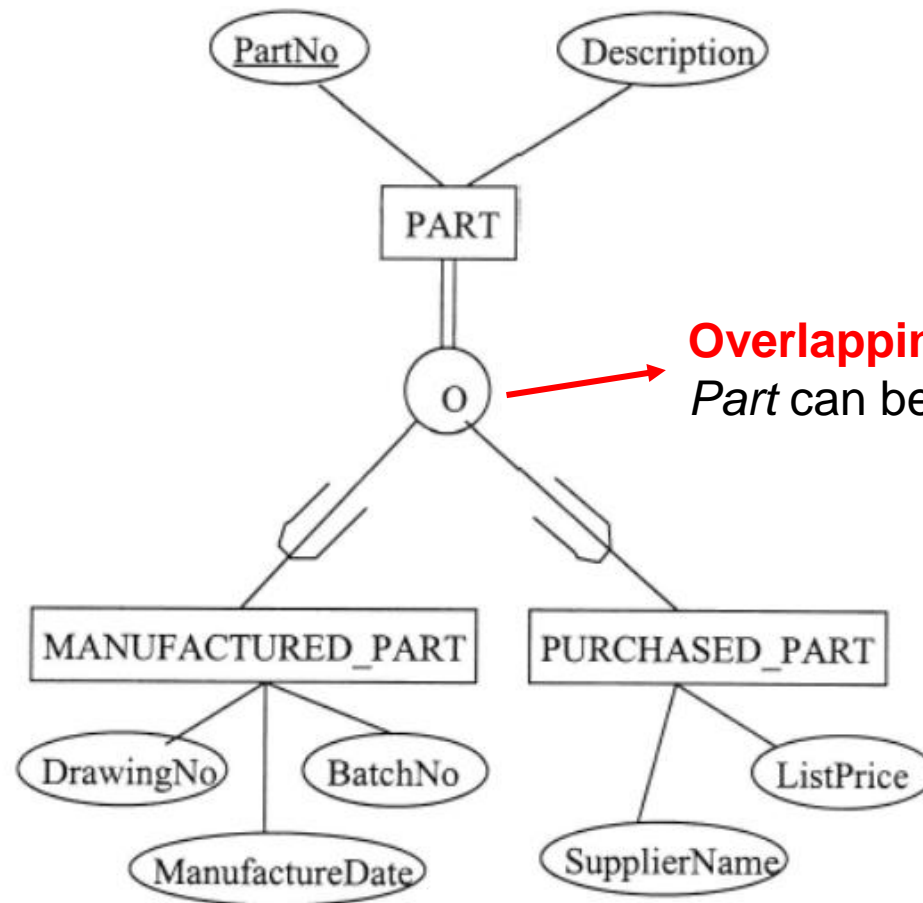
Generalization Hierarchy



Disjoint, a *Secretary* can not be a *Technician* or *Engineer* as well

Note: subset symbol on a line indicates direction of class/subclass relationship.

Generalization Hierarchy



Generalization Hierarchy

employee - generalized class contains only attributes in common

secretary, engineer, trucker - specializations of employee

•ATTRIBUTE INHERITANCE

an entity that is a member of a subclass inherits all the attributes of the entity as a member of the superclass. the entity will also inherit all *relationship* instances for relationship types in which the superclass participates.

•IMPLIED CONSTRAINTS

an entity cannot exist in database by being only a member of a subclass, it must also be member of superclass (reverse is not true).

Thank You.