

SQL Part 3

Database Modifications

Integrity constraints

Peter Scheuermann

Outline

- **Database Modifications**
- **Static Integrity Constraints**
 - ▶ Domain Constraints
 - ▶ Key / Referential Constraints
 - ▶ Semantic Integrity Constraints
- **Dynamic Integrity Constraints**
 - ▶ Triggers

Database Modifications

- A **modification** command does not return a result (as a query does), but changes the database in some way.
- Three kinds of modifications:
 1. **Insert** a tuple or tuples.
 2. **Delete** a tuple or tuples.
 3. **Update** the value(s) of an existing tuple or tuples.

Our Running Example

- Our SQL queries will be based on the following database schema.

- ▶ Underline indicates key attributes.

Beers(name, manf)

Bars(name, addr, license)

Drinkers(name, addr, phone)

Likes(drinker, beer)

Sells(bar, beer, price)

Frequents(drinker, bar)

Insertion

- To insert a single tuple:

```
INSERT INTO <relation>  
VALUES ( <list of values> );
```

- Example: add to Likes(drinker, beer) the fact that Sally likes Bud.

```
INSERT INTO Likes  
VALUES('Sally', 'Bud');
```

Specifying Attributes in INSERT

- We may add to the relation name a list of attributes.
- Two reasons to do so:
 1. We forget the standard order of attributes for the relation.
 2. We don't have values for all attributes, and we want the system to fill in missing components with NULL or a default value.

Example: Specifying Attributes

- Another way to add the fact that Sally likes Bud to `Likes(drinker, beer)`:

```
INSERT INTO Likes(beer, drinker)  
VALUES('Bud', 'Sally');
```

Adding Default Values

- In a CREATE TABLE statement, we can follow an attribute by DEFAULT and a value.
- When an inserted tuple has no value for that attribute, the default will be used.

Example: Default Values

```
CREATE TABLE Drinkers (  
    name CHAR(30) PRIMARY KEY,  
    addr CHAR(50)  
        DEFAULT '123 Sesame St.',  
    phone CHAR(16)  
);
```

Example: Default Values

```
INSERT INTO Drinkers(name)  
VALUES( 'Sally');
```

Resulting tuple:

name	address	phone
Sally	123 Sesame St	NULL

Inserting Many Tuples

- We may insert the entire result of a query into a relation, using the form:

```
INSERT INTO <relation>  
( <subquery> );
```

Example: Insert a Subquery

- Using `Frequents(drinker, bar)`, enter into the new relation `PotBuddies(name)` all of Sally's "potential buddies," i.e., those drinkers who frequent at least one bar that Sally also frequents.

Solution

The other
drinker

```
INSERT INTO PotBuddies
(SELECT d2.drinker
FROM Frequents d1, Frequents d2
WHERE d1.drinker = 'Sally' AND
d2.drinker <> 'Sally' AND
d1.bar = d2.bar
);
```

Pairs of Drinker
tuples where the
first is for Sally,
the second is for
someone else,
and the bars are
the same.

Deletion

- To delete tuples satisfying a condition from some relation:
DELETE FROM <relation>
WHERE <condition>;
- Important : FROM clause can specify only ONE relation !

Example: Deletion

- Delete from Likes(drinker, beer) the fact that Sally likes Bud:

DELETE FROM Likes

WHERE drinker = 'Sally' AND

beer = 'Bud';

Example: Delete all Tuples

- Make the relation Likes empty:

DELETE FROM *Likes*;

- Note no WHERE clause needed.

Example: Delete Some Tuples

- **Delete** from **Beers(name, manf)** all beers for which there is another beer by the same manufacturer.

DELETE FROM Beers b

WHERE EXISTS (

SELECT name FROM Beers

WHERE manf = b.manf AND

name <> b.name);

Beers with the same manufacturer and a different name from the name of the beer represented by tuple b.

Semantics of Deletion --- (1)

- Suppose Anheuser-Busch makes only Bud and Bud Lite.
- Suppose we come to the tuple b for Bud first.
- The subquery is nonempty, because of the Bud Lite tuple, so we delete Bud.
- Now, when b is the tuple for Bud Lite, do we delete that tuple too?

Semantics of Deletion --- (2)

- Answer: we *do* delete Bud Lite as well.
- The reason is that deletion proceeds in two stages:
 1. Mark all tuples for which the WHERE condition is satisfied.
 2. Delete the marked tuples.

Updates

- To change certain attributes in certain tuples of a relation:

UPDATE <relation>

SET <list of attribute assignments>

WHERE <condition on tuples>;

Example: Update

- Change drinker Fred's phone number to 555-1212:

```
UPDATE Drinkers  
SET    phone = '555-1212'  
WHERE  name = 'Fred';
```

Example: Update Several Tuples

- Make \$4 the maximum price for beer:

```
UPDATE Sells  
SET    price = 4.00  
WHERE  price > 4.00;
```

Semantic Integrity Constraints(ICs)

■ Objective:

- ▶ capture semantics of the mini-world in the database
- ▶ ensuring that authorized changes to the database do not result in a loss of **data consistency**
- ▶ guard against accidental damage to the database (avoid data entry errors)

■ Advantages of a centralized, automatic mechanism to ensures semantic integrity constraints:

- ▶ More effective integrity control
- ▶ Stored data is more faithful to real-world meaning
- ▶ Easier application development, better maintainability

■ Note: DBMS allow to capture more ICs than, e.g., ER Model

Integrity Constraint (IC)

- **Integrity Constraint (IC):**
condition that must be true for every instance of a database
 - ▶ A **legal** instance of a relation is one that satisfies all specified ICs
 - DBMS should never allow illegal instances....
- ICs are *specified* in the database schema
 - ▶ The database designer is responsible to ensure that the integrity constraints are not contradicting each other!
- ICs are *checked* when the database is modified
 - ▶ With one degree of freedom:
 - After a SQL statement, or at the end of a transaction?
- Possible *reactions* if an IC is violated:
 - ▶ Undoing of a database operation
 - ▶ Abort of the transaction
 - ▶ Execution of “maintenance” operations to make db legal again

Types of Integrity Constraints

■ Static Integrity Constraints

describe conditions that every *legal instance* of a database must satisfy

- ▶ Inserts / deletes / updates that violate ICs are disallowed
- ▶ Three kinds:
 - *Domain Constraints*
 - *Key Constraints & Referential Integrity*
 - *Semantic Integrity Constraints; Assertions*

■ Dynamic Integrity Constraints

are predicates on database state changes

- ▶ ***Triggers***

Domain (Attribute) Constraints

- The most elementary form of an integrity constraint:
- Fields must be of right data domain
 - ▶ always enforced for values inserted in the database
 - ▶ Also: queries are tested to ensure that the comparisons make sense.
- SQL DDL allows domains of attributes to be restricted in the **create table** definition with the following clauses:
 - ▶ **CHECK** (*condition on attribute*) *<- Simple semantic integrity constraint*
all values of the attribute must always satisfy the given predicate
 - ▶ **DEFAULT** *default-value*
default value for an attribute if its value is omitted in an insert stmt.
 - ▶ **NOT NULL**
attribute is not allowed to become NULL

Example of Domain Constraints

CREATE TABLE Student

```
(  
    sid      INTEGER      PRIMARY KEY,  
    name     VARCHAR(20)  NOT NULL,  
    semester INTEGER      DEFAULT 1 CHECK (semester > 0),  
    birthday DATE,  
    country  VARCHAR(20)  
);
```

Semantic:

sid is primary key of **Student**

name must not be NULL

semester will be 1 if not specified by an insert

all other attributes can be NULL (**birthday** and **country**)

Example:

```
INSERT INTO Student(sid,name) VALUES (123, 'Peter');
```

Example

```
CREATE TABLE Sells (  
    bar          CHAR(20) ,  
    beer         CHAR(20)   CHECK ( beer IN  
                                   (SELECT name FROM Beers) ) ,  
    price        REAL CHECK ( price <= 5.00 )  
);
```

Timing of Checks

- Attribute-based checks performed only when a value for that attribute is inserted or updated.
 - ▶ **Example:** **CHECK (price <= 5.00)** checks every new price and rejects the modification (for that tuple) if the price is more than \$5.
 - ▶ **Example:** **CHECK (beer IN (SELECT name FROM Beers))** not checked if a beer is deleted from Beers (unlike foreign-keys).

Tuple-Based Checks

- CHECK (<condition>) may be added as a relation-schema element.
- The condition may refer to any attribute of the relation.
 - ▶ But any other attributes or relations require a subquery.
- Checked on insert or update only.

Example: Tuple-Based Check

- Only Joe's Bar can sell beer for more than \$5:

```
CREATE TABLE Sells (  
    bar          CHAR(20) ,  
    beer         CHAR(20) ,  
    price        REAL,  
    CHECK (bar = 'Joe' 's Bar' OR  
           price <= 5.00)  
);
```

Semantic Integrity Constraints

- Integrity constraints on more than one attribute
- Also, a name for integrity constraint would be very useful for administration / maintenance...
- SQL:
CONSTRAINT *name* **CHECK** (*semantic-condition*)
- One can use subqueries to express constraint

Semantic Constraints Example

```
CREATE TABLE Enrollment
(
    sid      INTEGER      REFERENCES Student,
    c-code   VARCHAR(8)   REFERENCES Course ,
    empid    INTEGER      REFERENCES Lecturer,
    grade    INTEGER,
    CONSTRAINT grade CHECK (grade between 0 and 100),
    CONSTRAINT rightLecturer
        CHECK ( empid = (SELECT c.lecturer
                           FROM   Course c
                           WHERE  c.c_code=c-code) )
);
```

SQL: Naming Integrity Constraints

- The **CONSTRAINT** clause can be used to name all types of integrity constraints

- Example:

```
CREATE TABLE Enrolled
(
    sid          INTEGER,
    c-code       VARCHAR(8) ,
    grade        CHAR(2) ,
    CONSTRAINT FK_sid_enrolled    FOREIGN KEY (sid)
                                REFERENCES Student
                                ON DELETE CASCADE,
    CONSTRAINT FK_cid_enrolled    FOREIGN KEY (c-code)
                                REFERENCES Course
                                ON DELETE CASCADE,
    CONSTRAINT CK_grade_enrolled CHECK(grade in ('F' ,...)) ,
    CONSTRAINT PK_enrolled        PRIMARY KEY (sid, c-code)
);
```

Example: Deferring Constraints

```
CREATE TABLE Course
(
    c_code          VARCHAR(8) ,
    title           VARCHAR(220) ,
    lecturer        INTEGER,
    credit_points   INTEGER,
    CONSTRAINT Course_PK PRIMARY KEY (c_code) ,
    CONSTRAINT Course_FK FOREIGN KEY (lecturer)
        REFERENCES Lecturer DEFERRABLE INITIALLY DEFERRED
);
```

- Allows to insert a new course referencing a lecturer which is not present at that time, but who will be added later *in the same transaction*.
- Behaviour can be dynamically changed within a transaction with the SQL statement

SET CONSTRAINT Course_FK IMMEDIATE;

Deferring Constraint Checking

- Any constraint - domain, key, foreign-key, semantic - may be declared:
 - ▶ **NOT DEFERRABLE**
The default. It means that every time a database modification occurs, the constraint is checked immediately afterwards.
 - ▶ **DEFERRABLE**
Gives the option to wait until a transaction is complete before checking the constraint.
 - **INITIALLY DEFERRED** wait until transaction end, but allow to dynamically change later

Assertions

- The integrity constraints seen so far are associated with a single table
 - ▶ Plus: they are required to hold only if the associated table is nonempty!
- Need for more general integrity constraints
 - ▶ E.g. integrity constraints over several tables
 - ▶ Always checked, independent if one table is empty
- **Assertion**: a predicate expressing a condition that we wish the database always to satisfy.
- SQL-92 syntax:
create assertion *<assertion-name>* **check** (*<condition>*)
- Assertions are schema objects (like tables or views)
- When an assertion is made, the system tests it for validity, and tests it again on every update that may violate it
 - ▶ This testing may introduce a significant amount of overhead; hence assertions should be used with great care.

Assertion Example 1

- Asserting $\forall X : P(X)$ is achieved in a round-about fashion using $\text{not exists } X \text{ such that not } P(X)$
- Example: For all students, the sum of all grades for a course must be less or equal than 10000

```
CREATE ASSERTION grade-constraint CHECK  
(  
  not exists ( select   c-code  
                from     Enrollment  
                group by c_code  
                having    sum(grade) > 10000 )  
)
```

Example Assertion 2

- In **Drinkers(name, addr, phone)** and **Bars(name, addr, license)**, there cannot be more bars than drinkers.

```
CREATE ASSERTION FewBar CHECK (  
    (SELECT COUNT (*) FROM Bars) <=  
    (SELECT COUNT (*) FROM Drinkers)  
);
```

Timing of Assertion Checks

- In principle, we must check every assertion after every modification to any relation of the database.
- A clever system can observe that only certain changes could cause a given assertion to be violated.
 - ▶ **Example:** No change to Beers can affect FewBar. Neither can an insertion to Drinkers.