

Transactions

Peter Scheuermann

Transaction Concept

- Many enterprises and organizations use databases to store information about their state
 - ▶ e.g., Balances of all depositors at a bank
- When an event occurs in the real world that changes the state of the enterprise, a program is executed to change the database state in a corresponding way
 - ▶ e.g., Bank balance must be updated when deposit is made
- Such a program is called a **transaction**:
a collection of one or more operations on one or more databases, which reflects a discrete unit of work
 - ▶ In the real world, this happened (completely) or it didn't happen at all (Atomicity)

What Does a Transaction Do?

■ Return information from the database

- ▶ RequestBalance transaction:
Read customer's balance in database and output it
=> transactions can be read-only

■ *Update the database to reflect the occurrence of a real world event*

- ▶ *Transfer* money between accounts
 - Update customers' balances in database(s)
- ▶ Purchase a group of products
- ▶ Students enrolling in an unit of study

■ Cause the occurrence of a real world event

- ▶ Withdraw transaction:
Dispense cash (and update customer's balance in database)

Transaction Concept (continued)

- A **transaction** is a *unit* of program execution that accesses and possibly updates various data items.

- E.g. transaction to transfer \$50 from account A to account B:

1. **read**(A)
2. $A := A - 50$
3. **write**(A)
4. **read**(B)
5. $B := B + 50$
6. **write**(B)

- Two main issues to deal with:

- ▶ Failures of various kinds, such as hardware failures and system crashes
- ▶ Concurrent execution of multiple transactions

Example of Fund Transfer

- Transaction to transfer \$50 from account A to account B:
 1. **read**(A)
 2. $A := A - 50$
 3. **write**(A)
 4. **read**(B)
 5. $B := B + 50$
 6. **write**(B)
- **Atomicity requirement**
 - ▶ if the transaction fails after step 3 and before step 6, money will be “lost” leading to an inconsistent database state
 - Failure could be due to software or hardware
 - ▶ the system should ensure that updates of a partially executed transaction are not reflected in the database
- **Durability requirement** — once the user has been notified that the transaction has completed (i.e., the transfer of the \$50 has taken place), the updates to the database by the transaction must persist even if there are software or hardware failures.

Example of Fund Transfer (Cont.)

- Transaction to transfer \$50 from account A to account B:
 1. **read**(A)
 2. $A := A - 50$
 3. **write**(A)
 4. **read**(B)
 5. $B := B + 50$
 6. **write**(B)
- **Consistency requirement** in above example:
 - ▶ the sum of A and B is unchanged by the execution of the transaction
- In general, consistency requirements include
 - **Explicitly specified integrity constraints** such as primary keys and foreign keys
 - **Implicit integrity constraints**
 - e.g. sum of balances of all accounts, minus sum of loan amounts must equal value of cash-in-hand
 - ▶ A transaction must see a consistent database.
 - ▶ During transaction execution the database may be temporarily inconsistent.
 - ▶ **When the transaction completes successfully the database must be consistent**
 - **Erroneous transaction logic can lead to inconsistency**

Example of Fund Transfer (Cont.)

- **Isolation requirement** — if between steps 3 and 6, another transaction T2 is allowed to access the partially updated database, it will see an inconsistent database (the sum $A + B$ will be less than it should be).

T1

1. **read**(A)
2. $A := A - 50$
3. **write**(A)
4. **read**(B)
5. $B := B + 50$
6. **write**(B)

T2

read(A), read(B), print(A+B)

- **Isolation can be ensured trivially by running transactions serially**
 - ▶ that is, one after the other.
- **However, executing multiple transactions concurrently has significant benefits, as we will see later.**

Transactional Guarantees

- The execution of each transaction must maintain relationship between the database state and the enterprise state
 - ▶ correctness and consistency of the database is paramount!
- Therefore additional requirements are placed on the execution of transactions beyond those placed on ordinary programs:

- ▶ Atomicity
- ▶ Consistency
- ▶ Isolation
- ▶ Durability

**ACID
properties**

A C I D Properties


- **Atomicity.** Transaction should either complete or have no effect at all
 - ▶ In case of a failure, all effects of operations of not-completed transactions are undone.
- **Consistency.** Execution of a transaction in isolation preserves the consistency of the database.
- **Isolation.** Although multiple transactions may execute concurrently, each transaction must be unaware of other concurrently executing transactions.
 - ▶ Intermediate transaction results must be hidden from other concurrently executed transactions.
- **Durability.** The effect of a transaction on the database state should not be lost once the transaction has committed

ACID properties handled transparently for the transaction by the DBMS

Transaction State

- **Active** – the initial state; the transaction stays in this state while it is executing
- **Partially committed** – after the final statement has been executed.
- **Failed** -- after the discovery that normal execution can no longer proceed.
- **Aborted** – after the transaction has been rolled back and the database restored to its state prior to the start of the transaction. Two options after it has been aborted:
 - ▶ **restart the transaction**
 - can be done only if no internal logical error
 - ▶ **kill the transaction**
- **Committed** – after successful completion.

Transaction State (Cont.)



A - Atomicity

- A real-world event either happens or does not happen
 - ▶ Student either registers or does not register
- Similarly, the system must ensure that either the corresponding transaction runs to completion or, if not, it has no effect at all
 - ▶ a user can think of a transaction as always executing all its actions in one step, or not executing any actions at all.
 - ▶ Not true of ordinary programs. A crash could leave files partially updated on recovery.
 - DBMS *logs* all actions so that it can *undo* the actions of aborted transactions.
 - Also, in case of a failure, all actions of not-committed transactions are *undone*.

API for Transactions

- Data manipulation language must provide commands for setting transaction boundaries. For example:
 - ▶ begin transaction
 - ▶ commit ; rollback
- In many DBMS such as Oracle, a transaction begins implicitly
 - ▶ Some other DBMS (eg. SQL Server or PostgreSQL) provide a **BEGIN TRANSACTION** command
- A transaction ends by:
 - ▶ **COMMIT** *requests* to **commit** current transaction
 - The system might commit the transaction, or it might abort if needed.
 - ▶ **ROLLBACK** causes current transaction to **abort** - *always satisfied*.
- The commit command is a request
 - ▶ The system might commit the transaction, or it might abort it

Transaction Example

- Pseudocode for a product order transaction:

display greeting

get order request

BEGIN TRANSACTION

SELECT *product record*

IF product is available THEN

UPDATE quantityOnOrder of product record

INSERT order record

send message to shipping departement

COMMIT

ELSE

ROLLBACK

END IF

Another Transaction Example

■ Transaction in Embedded SQL

```
1. EXEC SQL BEGIN DECLARE SECTION
2.     int flight;
3.     char date[10]
4.     char seat [3]
5.     int occ;
6. EXEC SQL END DECLARE SECTION
7. START TRANSACTION
8. Void chooseSeat() {
9.     /* C code to prompt the user to enter flight, date, and seat and store these in the three variables
       with those name */
10. EXEC SQL Select occupied into :occ
11.         From Flights
12.         Where fltNum=:flight and fltDate=:date and fltSeat=:seat;
13. If (!occ) {
14.     EXEC SQL Update Flights
15.         Set occupied = true
16.         Where fltNum=:flight and fltDate=:date and fltSeat=:seat;
17. /*C and SQL code to record the seat assignment and inform the user of the assignment */
18. }
19. Else /* C code to notify user of unavailability and ask for another seat selection */
20. } EXEC COMMIT;
```

Database Consistency

- **Enterprise (Business) Rules** limit the occurrence of certain real-world events
 - ▶ Student cannot register for a course if the current number of registrants equals the maximum allowed
- **Correspondingly, allowable database states are restricted**
 - ▶ $cur_reg \leq max_reg$
- These limitations are called (static) **integrity constraints**: assertions that must be satisfied by the database state
- **Database is consistent** if all static integrity constraints are satisfied

Transaction Consistency

- A consistent database state does not necessarily model the actual state of the enterprise
 - ▶ A deposit transaction that increments the balance by the wrong amount maintains the integrity constraint $balance \geq 0$, but does not maintain the relation between the enterprise and database states
 - ▶ **Dynamic Integrity Constraints:** Some constraints restrict allowable state transitions
 - A transaction might transform the database from one consistent state to another, but the transition might not be permissible
 - **Example:** Students can only progress from Junior to the Senior year, but can never be degraded.
- A consistent transaction maintains database consistency *and* the correspondence between the database state and the enterprise state (implements its specification)
 - ▶ Specification of deposit transaction includes
$$balance = balance' + amt_deposit ,$$
$$(balance' \text{ is the initial value of } balance)$$

Transaction Consistency (cont'd)

- A **transaction is consistent** if, assuming the database is in a consistent state initially, when the transaction completes:
 - ▶ **All static integrity constraints are satisfied** (but constraints might be violated in intermediate states)
 - Can be checked by examining snapshot of database
 - ▶ **New state satisfies specifications of transaction**
 - Cannot be checked from database snapshot
 - ▶ **No dynamic constraints have been violated**
 - Cannot be checked from database snapshot

Integrity Constraints and Transactions

■ When do we check integrity constraints?

- ▶ *Immediate* after an SQL statement or *at the end* of a transaction?

■ Remember from last week:

- ▶ Integrity constraints may be declared:

- **NOT DEFERRABLE**

The default. It means that every time a database modification occurs, the constraint is checked immediately afterwards.

- **DEFERRABLE**

Gives the option to wait until a transaction is complete before checking the constraint.

Deferrable Integrity Constraints Example

```
CREATE TABLE COURSE (  
    c_code          VARCHAR(8) ,  
    title           VARCHAR(220) ,  
    lecturer        INTEGER ,  
    credit_points   INTEGER ,  
    CONSTRAINT COURSE_PK PRIMARY KEY (c_code) ,  
    CONSTRAINT COURSE_FK FOREIGN KEY (lecturer)  
        REFERENCES Lecturer DEFERRABLE INITIALLY IMMEDIATE  
); //this is Oracle syntaz  
  
;  
  
INSERT INTO Course VALUES ('EECS339' , 'DBMS' , 42 , 1) ;  
INSERT INTO Lecturer VALUES (42 , 'Charley Sheen' , ...) ;  
COMMIT ;
```

lecturer 42 has
to exist for the
FK to be OK

D - Durability

- The system must ensure that once a transaction commits, its effect on the database state is not lost in spite of subsequent failures
 - ▶ Not true of ordinary programs. A media failure after a program successfully terminates could cause the file system to be restored to a state that preceded the program's execution
- Implementing Durability:
 - ▶ Database is stored redundantly on mass storage devices to protect against media failure
 - ▶ Architecture of mass storage devices affects type of media failures that can be tolerated
 - ▶ Related to **Availability**: extent to which a (possibly distributed) system can provide service despite failure
 - Non-stop DBMS (mirrored disks)
 - Recovery based DBMS (log)

I - Isolation

■ **Serial Execution: transactions execute in sequence**

- ▶ Each one starts after the previous one completes.
 - Execution of one transaction is not affected by the operations of another since they do not overlap in time
- ▶ The execution of each transaction is **isolated** from all others.

■ If the initial database state and all transactions are consistent, then the final database state will be consistent and will accurately reflect the real-world state, *but*

- ▶ Serial execution is inadequate from a performance perspective

■ **Concurrent execution offers performance benefits:**

- ▶ A computer system has multiple resources capable of executing independently (e.g., cpu's, I/O devices), *but*
- ▶ A transaction typically uses only one resource at a time
- ▶ *but* might not be correct...

Schedules

- **Schedule** – a sequences of instructions that specify the chronological order in which instructions of concurrent transactions are executed
 - ▶ a schedule for a set of transactions must consist of all instructions of those transactions
 - ▶ must preserve the order in which the instructions appear in each individual transaction.
- A transaction that successfully completes its execution will have a commit instructions as the last statement
 - ▶ by default transaction assumed to execute commit instruction as its last step
- A transaction that fails to successfully complete its execution will have an abort instruction as the last statement

Schedule 1

- Let T_1 transfer \$50 from A to B , and T_2 transfer 10% of the balance from A to B .
- **A serial schedule in which T_1 is followed by T_2 :**

T_1	T_2
read (A) $A := A - 50$ write (A) read (B) $B := B + 50$ write (B) commit	read (A) $temp := A * 0.1$ $A := A - temp$ write (A) read (B) $B := B + temp$ write (B) commit

Schedule 2

- A serial schedule where T_2 is followed by T_1

T_1	T_2
read (A) $A := A - 50$ write (A) read (B) $B := B + 50$ write (B) commit	read (A) $temp := A * 0.1$ $A := A - temp$ write (A) read (B) $B := B + temp$ write (B) commit

Schedule 3

- Let T_1 and T_2 be the transactions defined previously.
The following schedule is not a serial schedule, but it is *equivalent* to Schedule 1.

T_1	T_2
read (A) $A := A - 50$ write (A)	read (A) $temp := A * 0.1$ $A := A - temp$ write (A)
read (B) $B := B + 50$ write (B) commit	read (B) $B := B + temp$ write (B) commit

In Schedules 1, 2 and 3, the sum $A + B$ is preserved.

Schedule 4

- The following concurrent schedule does not preserve the value of $(A + B)$.

T_1	T_2
read (A) $A := A - 50$	read (A) $temp := A * 0.1$ $A := A - temp$ write (A) read (B)
write (A) read (B) $B := B + 50$ write (B) commit	$B := B + temp$ write (B) commit

Serializability

- **The Issue:** Maintaining database correctness when many transactions are accessing the database concurrently
 - ▶ Basic Assumption: *Each transaction preserves database consistency.*
 - ▶ Thus serial execution of a set of transactions preserves database consistency.
- **Serializability:**

A schedule is **serializable** if it is equivalent to a serial schedule

 - ▶ Different notions of schedule equivalence...
 - ▶ This lecture concentrates on **conflict serializability**

Simplified view of transactions

- ▶ We ignore operations other than **read** and **write** instructions
- ▶ We assume that transactions may perform arbitrary computations on data in local buffers in between reads and writes.
- ▶ Our simplified schedules consist of only **read** and **write** instructions.

Conflict Serializability

- Two schedules are *conflict serializable* if:
 - ▶ They involve the same actions of the same transactions
 - ▶ Every pair of conflicting actions is ordered the same way
- Two actions a_i and a_j of transactions T_i and T_j **conflict** if and only if they access the same data X and at least one of these actions wrote X . (a_i, a_j) are called a **conflict pair**.
 1. $a_i = \text{read}(X)$, $a_j = \text{read}(X)$. don't conflict.
 2. $a_i = \text{read}(X)$, $a_j = \text{write}(X)$. they conflict.
 3. $a_i = \text{write}(X)$, $a_j = \text{read}(X)$. they conflict
 4. $a_i = \text{write}(X)$, $a_j = \text{write}(X)$. they conflict
- **Note:** With SQL - Select corresponds to read, Insert, Delete, Update correspond to write

Conflict Serializability (Cont.)

- **Schedule 3 can be transformed into Schedule 6**, a serial schedule where T_2 follows T_1 , by series of swaps of non-conflicting instructions. Therefore Schedule 3 is conflict serializable.

T_1	T_2
read (A) write (A)	read (A) write (A)
read (B) write (B)	read (B) write (B)

Schedule 3

T_1	T_2
read (A) write (A) read (B) write (B)	read (A) write (A) read (B) write (B)

Schedule 6

Conflict Serializability (Cont.)

- Example of a schedule that is not conflict serializable:

T_3	T_4
read (Q)	
write (Q)	write (Q)

- We are unable to swap instructions in the above schedule to obtain either the serial schedule $\langle T_3, T_4 \rangle$, or the serial schedule $\langle T_4, T_3 \rangle$.

Serializable Schedule Example 1

- The concurrent schedule

$S: r_1(x) w_2(z) w_1(y)$

is equivalent to the serial schedules of T_1 and T_2 in either order:

- ▶ T1, T2: $r_1(x) w_1(y) w_2(z)$ and
- ▶ T2, T1: $w_2(z) r_1(x) w_1(y)$

- Reason: operations of distinct transactions on different data items commute.
- Hence, S is a *serializable* schedule

Serializable Schedule Example 2

- The concurrent schedule

$S: r1(z) r2(q) w2(z) r1(q) w1(y)$

is equivalent to the serial schedule $T1, T2$:

$r1(z) r1(q) w1(y) r2(q) w2(z)$

since read operations of distinct transactions on the same data item commute.

- Hence, S is a serializable schedule
- However, S is not equivalent to $T2, T1$ since read and write operations (or two write operations) of distinct transactions on the same data item do not commute.

Non-Serializable Schedule Example

- Example: course registration; *cur_reg* is the number of current registrants

$T1: r(cur_reg : 29) \qquad \qquad \qquad w(cur_reg : 30)$
 $T2: \qquad \qquad r(cur_reg : 29) \ w(cur_reg : 30)$

- Schedule not equivalent to $T1, T2$ or $T2, T1$
- Database state no longer corresponds to real-world state, integrity constraint violated


Testing for Conflict Serializability

- Consider some schedule of a set of transactions T_1, T_2, \dots, T_n

- **Precedence graph:**

- ▶ direct graph where the vertices are the transactions.
- ▶ edge from T_i to T_j if the two transaction conflict, and T_i accessed the data item on which the conflict arose earlier.

- Example:
 T_1 and T_2 have
2 conflict pairs



- Central Theorem:
A schedule is **conflict serializable** if and only if its *precedence graph is acyclic*.
 - ▶ The serializability order can be obtained by a *topological sorting* of the graph.