

SQL Triggers

Peter Scheuermann

Outline

■ Static Integrity Constraints

- ▶ Domain Constraints
- ▶ Key / Referential Constraints
- ▶ Semantic Integrity Constraints

■ Dynamic Integrity Constraints

Triggers

- A **trigger** is a statement that is executed automatically if specified modifications occur to the DBMS.
- A trigger specification consists of three parts:
ON event IF condition THEN action
 - ▶ *Event* (what activates the trigger?)
 - ▶ *Condition* (guard / test whether the trigger shall be executed)
 - ▶ *Action* (what happens if the trigger is run)
- Also called **event-condition (ECA)** rules
- Triggers introduced to SQL standard in SQL:1999, but supported even earlier using non-standard syntax by most databases.

Why Triggers?

■ Constraint maintenance

- ▶ Triggers can be used to maintain foreign-key and semantic constraints; commonly used with **ON DELETE** and **ON UPDATE**

■ Business rules

- ▶ Some dynamic business rules can be encoded as triggers

■ Monitoring

- ▶ E.g. to react on the insertion of some kind of sensor reading into db

■ Maintenance of auxiliary cached data

- ▶ Careful! Many systems now support *materialized views* which should be preferred against such maintenance triggers

■ Simplified application design

- ▶ E.g. exceptions modelled as update operations on a database (if applicable)

Trigger example

```
CREATE TRIGGER CPS116AutoRecruit
```

```
AFTER INSERT ON Student → Event
```

```
REFERENCING NEW ROW AS newStudent
```

```
FOR EACH ROW
```

```
WHEN (newStudent.GPA > 3.0) → Condition
```

```
INSERT INTO Enroll
```

```
VALUES (newStudent.SID, 'CPS116');
```

↓
Action

Trigger options

■ Possible events include:

- ▶ *INSERT ON table*
- ▶ *DELETE ON table*
- ▶ *UPDATE [OF column] ON table*

■ Granularity—trigger can be activated:

- ▶ *FOR EACH ROW* modified
- ▶ *FOR EACH STATEMENT* that performs modification

■ Timing—action can be executed:

- ▶ *AFTER* or *BEFORE* the triggering event

Trigger Example (SQL:1999)

```
CREATE TRIGGER gradeUpgrade  
  AFTER INSERT INTO Assessment  
  REFERENCING NEW TABLE Assess  
  FOR EACH STATEMENT  
    UPDATE Enrolled  
      SET grade='P'  
    WHERE Enrolled.sid=Assess.sid AND  
      Enrolled.ucode=Assess.ucode AND  
      Assess.mark >= 50
```

Triggering Events and Actions in SQL

- Triggering event can be **insert, delete or update**
- Triggers on update can be restricted to specific attributes
`CREATE TRIGGER overdraft-trigger AFTER UPDATE OF balance
ON account`
- Values of attributes before and after an update can be referenced
 - ▶ **REFERENCING OLD ROW AS name** : for deletes and updates
 - ▶ **REFERENCING NEW ROW AS name** : for inserts and updates
- Triggers can be activated before an event, which can serve as extra constraints.
 - ▶ E.g. convert blanks to null:
`CREATE TRIGGER Setnull-trigger BEFORE UPDATE ON S
REFERENCING NEW ROW AS nrow
FOR EACH ROW
WHEN nrow.country = ' '
SET nrow.country = null`

Transition variables

- **OLD ROW**: the modified row before the triggering event
- **NEW ROW**: the modified row after the triggering event
- **OLD TABLE**: a hypothetical read-only table containing all modified rows before the triggering event
- **NEW TABLE**: a hypothetical table containing all modified rows after the triggering event

☞ **Not all of them make sense all the time**, e.g.

- ▶ *AFTER INSERT* statement-level triggers

- Can use only `NEW TABLE`

- ▶ *BEFORE DELETE* row-level triggers

- Can use only `OLD ROW`

- ▶ etc.

Trigger Granularity

■ Granularity

- ▶ *Row-level granularity*: change of a single row is an event (a single UPDATE statement might result in multiple events)
- ▶ *Statement-level granularity*: events are statements (a single UPDATE statement that changes multiple rows is a single event).

- Can be more efficient when dealing with SQL statements that update a large number of rows...

After Trigger Example

(statement granularity)

Keep track of salary averages in the log

```
CREATE TRIGGER RecordNewAverage
AFTER UPDATE OF Salary ON Employee
FOR EACH STATEMENT
BEGIN
    INSERT INTO Log
    VALUES (CURRENT_DATE, SELECT AVG(Salary)
                                FROM Employee );
END;
```

Trigger Granularity - Syntax

- Instead of executing a separate action for each affected row, a single action can be executed for all rows affected by a transaction
 - ▶ Use **FOR EACH STATEMENT** instead of **for each row**
 - ▶ Use **REFERENCING OLD TABLE**
or **REFERENCING NEW TABLE**
to refer to temporary tables (called *transition tables*) containing the affected rows

Triggers in SQL:1999

- **Events:** INSERT, DELETE, or UPDATE statements or changes to individual rows caused by these statements
 - ▶ Since SQL:2008: also **INSTEAD OF** triggers
- **Condition:** Anything that is allowed in a WHERE clause
- **Action:** An individual SQL statement or a program written in the language of **Procedural Stored Modules (PSM)** (which can contain embedded SQL statements)

Before Trigger Example (row granularity)

```
CREATE TRIGGER Max_EnrollCheck
  BEFORE INSERT ON Transcript
  REFERENCING NEW AS N                --row to be added
  FOR EACH ROW
  WHEN ( (SELECT COUNT (T.studId)
          FROM Transcript T
          WHERE T.CrsCode = N.CrsCode AND
                T.semester = N.semester)
        >=
        (SELECT CRSLIMIT.maxEnroll
          FROM CRSLIMIT U
          WHERE U.CrsCode = N.CrsCode And
                U.semester =N.semester)
        ) )
  BEGIN
    ROLLBACK;
  END;
```

*Check that
enrollment \leq limit*

Some Tips on Triggers

- Use BEFORE triggers
 - ▶ For checking integrity constraints
- Use AFTER triggers
 - ▶ For integrity maintenance and update propagation
- In Oracle, triggers cannot access “mutating” tables
 - ▶ e.g. AFTER trigger on the same table which just updates

Triggers in SQL:1999

- **Consideration:** *Immediate*

- ▶ Condition can refer to both the state of the affected row or table before *and* after the event occurs

- **Execution:** *Immediate* – can be before or after the execution of the triggering event

- ▶ Action of before trigger cannot modify the database

- **Granularity:** Both *row-level* and *statement-level*

Design Space of Triggers

- **Activation** - Occurrence of the *event*
- **Consideration** - The point, after activation, when *condition* is evaluated
 - ▶ Immediate or deferred (when the transaction requests to commit)
 - ▶ *Condition* might refer to both the state before and the state after *event* occurs
- **Execution** – point at which *action* occurs
 - ▶ With deferred consideration, execution is also deferred
 - ▶ With immediate consideration, execution can occur immediately after consideration or it can be deferred
 - If execution is immediate, execution can occur before, after, or instead of triggering event.
 - Before triggers adapt naturally to maintaining integrity constraints: violation results in rejection of event.

You should now be able to:

- **Capture Integrity Constraints in an SQL Schema**
 - ▶ Including key constraints, referential integrity, domain constraints and semantic constraints
 - ▶ And simple triggers for dynamic constraints
- Formulate complex semantic constraints using Assertions
- Know when to use Assertions, when triggers, and when CHECK constraints
- Know the semantic of deferring integrity constraints
- **Be able to formulate simple triggers**
- Know the difference between row-level & statement-level triggers